Exploring The Possibilities of Azure Fabric Abuses

Viktor Gazdag | Cloud Village | DefCon 33

WhoAmI

- Principal Security Consultant NCC Group
- Cloud and CI/CD Security
- Azure Certs: AZ-104, AZ-400, AZ-500, SC-100, CARTP, CAWASP
- Black Hat USA 2022, DefCon 31/32/33 Cloud Village, DoD CyberDT XSWG #15, DefCamp

What Is Fabric?

- “An enterprise-ready, end-to-end analytics platform. It unifies data movement, data processing, ingestion, transformation, real-time event routing, and report building. It supports these capabilities with integrated services like Data Engineering, Data Factory, Data Science, Real-Time Intelligence, Data Warehouse, and Databases.”

- SaaS analytics platform including Data Engineering, Data Warehouse, Synapse, Database etc.

https://github.com/woodspeed/conferences

Tenant Level Settings

- The Top 3 Settings To Review Are Related To:
- Sharing
- Service Principal
- User Function

Tenant Level

- Export and Sharing
  - External Data Sharing
  - Guest Users Can Access Microsoft Fabric

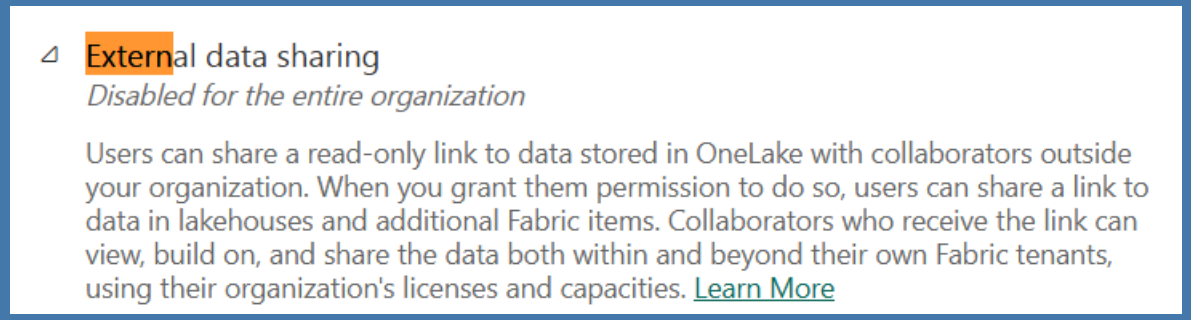

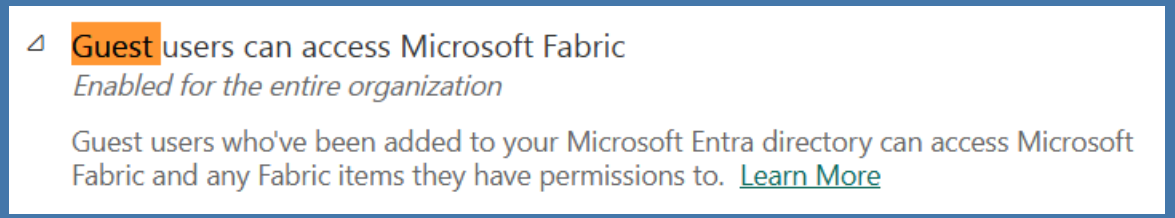

Tenant Level

- Service Principal Can Access Read-only Admin APIs
- Service Principal Can Call Fabric Public APIs

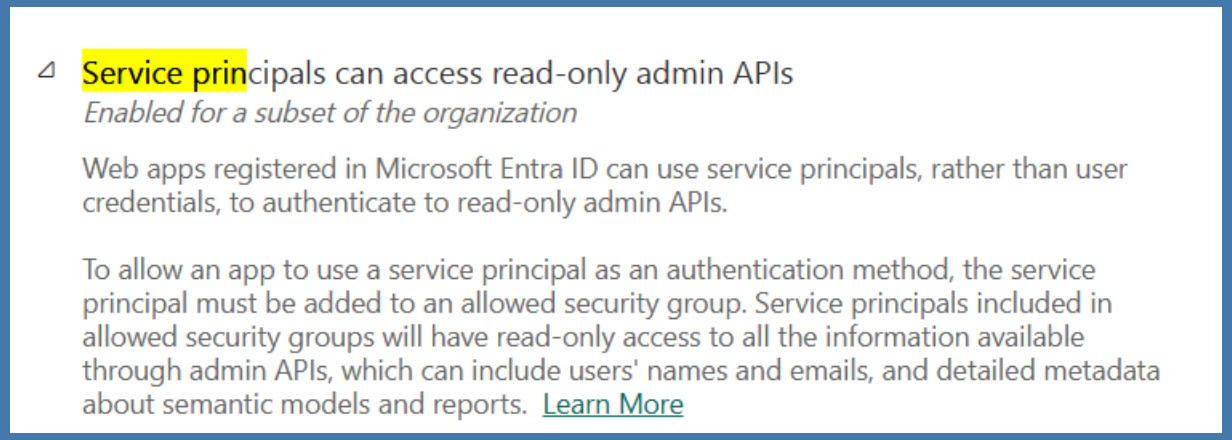

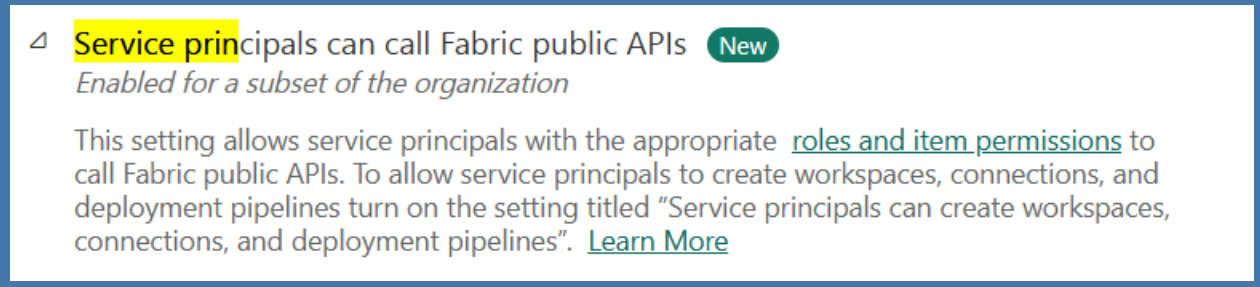

Tenant Level

- Restrict Content With Protected Labels From Being Shared Via Link With Everyone In Your Organization

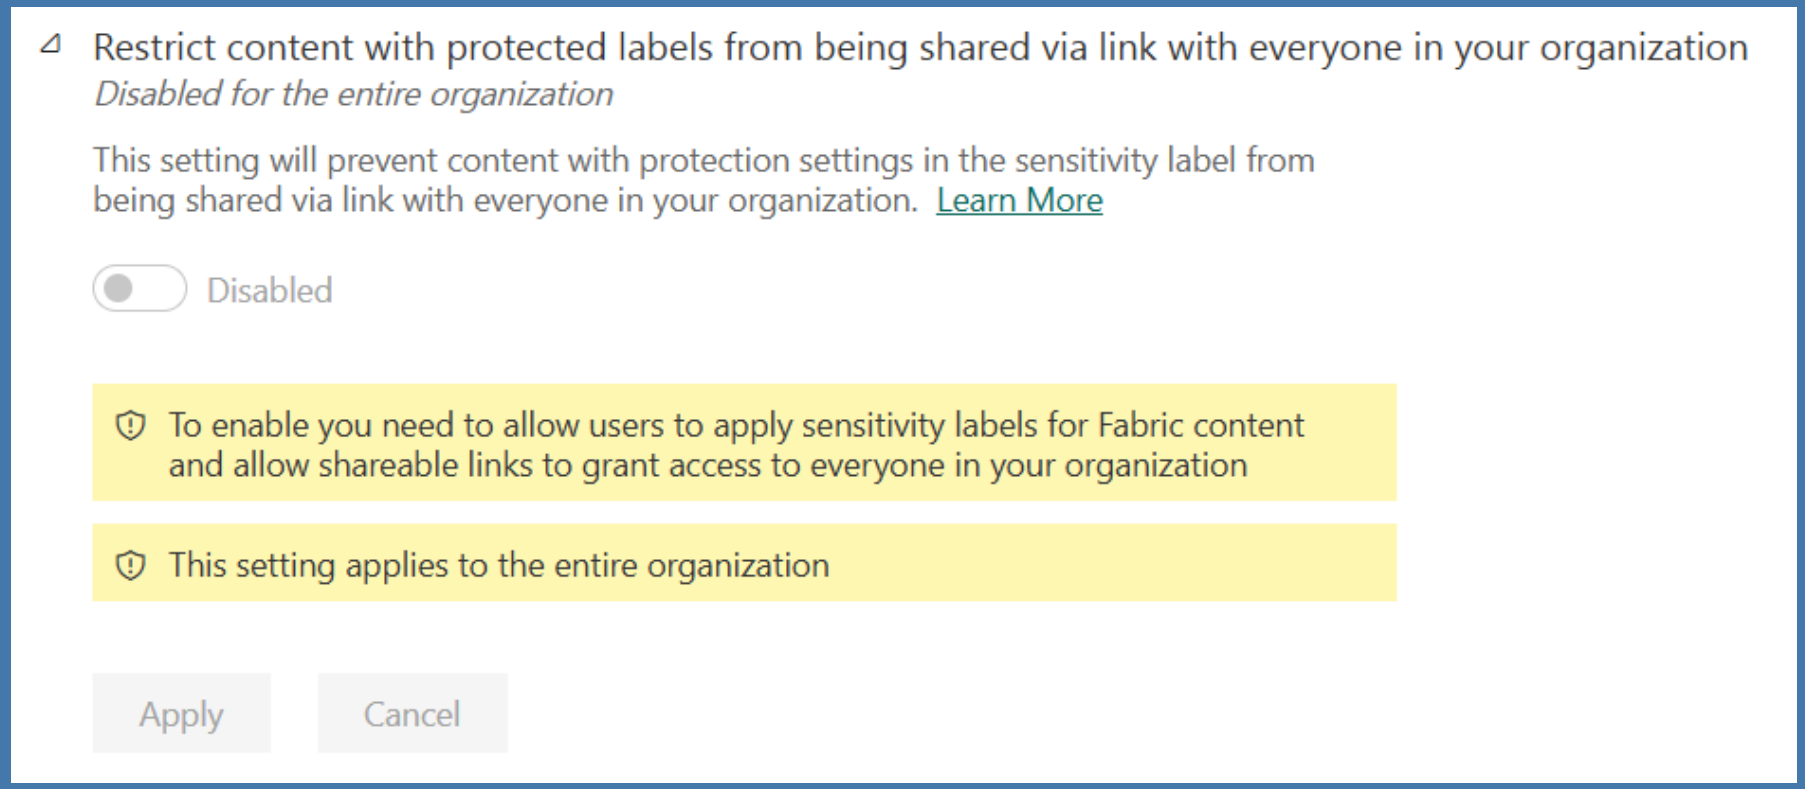

Tenant Level

- User Data Functions (Only To Specific Groups)

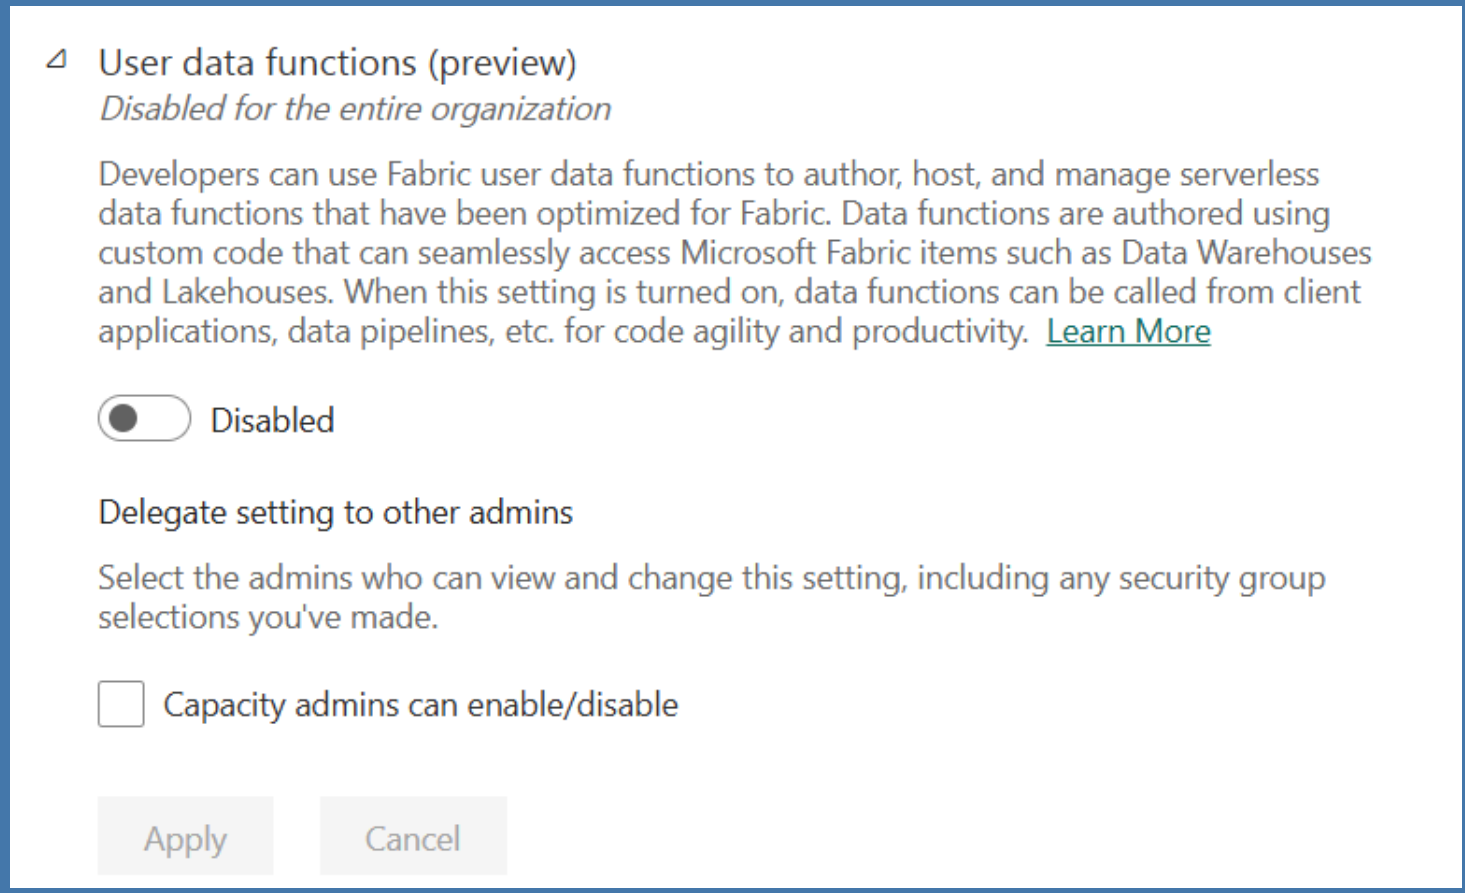

Tenant Level

- There Are More Than 3 Settings
- These Were Only The Top 3 Settings I Pick

- Conditional Access Policy
- Private Links For Secure Access To Fabric

Exploring Possibilities

- Backdoor
- Data Exfiltration

Backdoor

- Activator (Can Call Below)
- Notebook (Python SDK, 3rd Party Library)
- Power Automate (Low Code No Code, Entra ID Steps)
- Scheduled Spark Job (Python Code, Library Management Via Environment)

Backdoor

- Activator
  - Event Trigger
- Notebook
  - Azure Python SDK
  - Library Management
- User Or Identity
  - User
  - Enable Identity At Tenant Level Settings

Backdoor

- What Is Activator
  - Event Triggers For Fabric
    - Job Events
    - OneLake Events
    - Workspace Item Events
    - Azure Blog Storage Events

- Create/Update/Delete
- Succeeded/Failed
- Delay

Backdoor

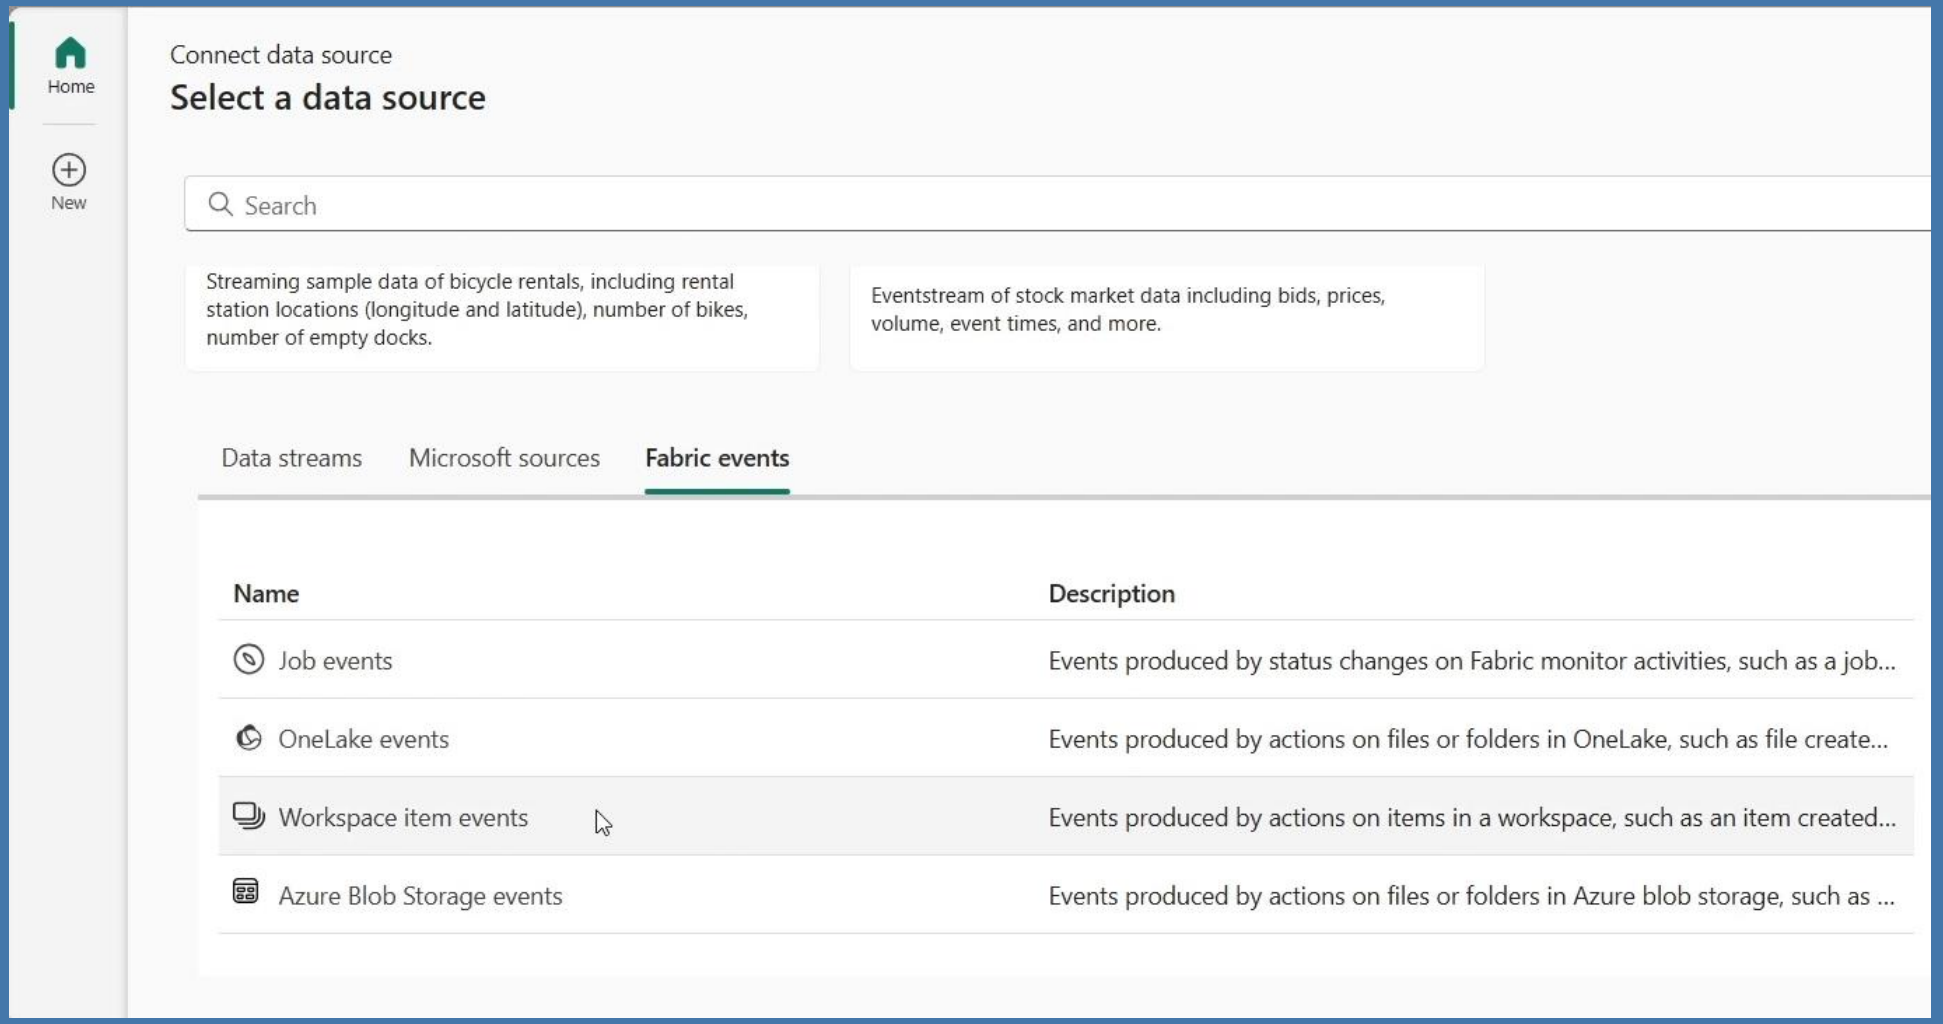
Home
New
Connect data source
Select a data source
Search
Streaming sample data of bicycle rentals, including rental station locations (longitude and latitude), number of bikes, number of empty docks.
Eventstream of stock market data including bids, prices, volume, event times, and more.
Data streams
Microsoft sources
Fabric events
Name
Description
Job events
Events produced by status changes on Fabric monitor activities, such as a job...
OneLake events
Events produced by actions on files or folders in OneLake, such as file create...
Workspace item events
Events produced by actions on items in a workspace, such as an item created...
Azure Blob Storage events
Events produced by actions on files or folders in Azure blob storage, such as ...

Backdoor

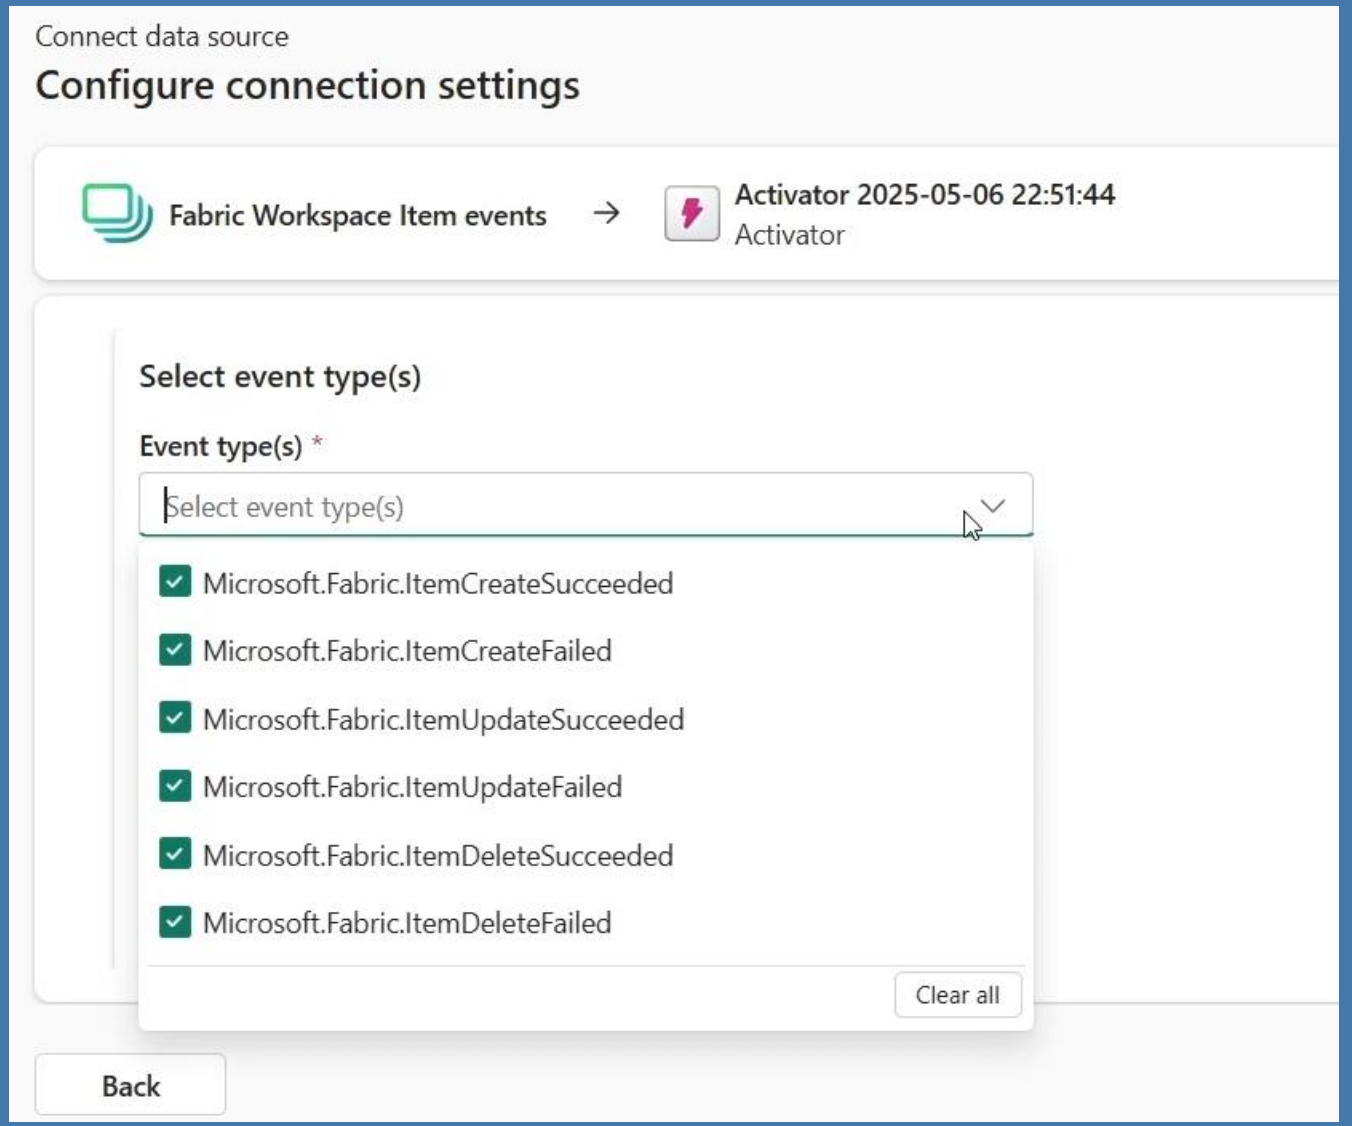
Connect data source
Configure connection settings
Fabric Workspace Item events
Activator 2025-05-06 22:51:44
Activator
Select event type(s)
Event type(s) *
Select event type(s)
Microsoft.Fabric.ItemCreateSucceeded
Microsoft.Fabric.ItemCreateFailed
Microsoft.Fabric.ItemUpdateSucceeded
Microsoft.Fabric.ItemUpdateFailed
Microsoft.Fabric.ItemDeleteSucceeded
Microsoft.Fabric.ItemDeleteFailed
Clear all
Back

Backdoor

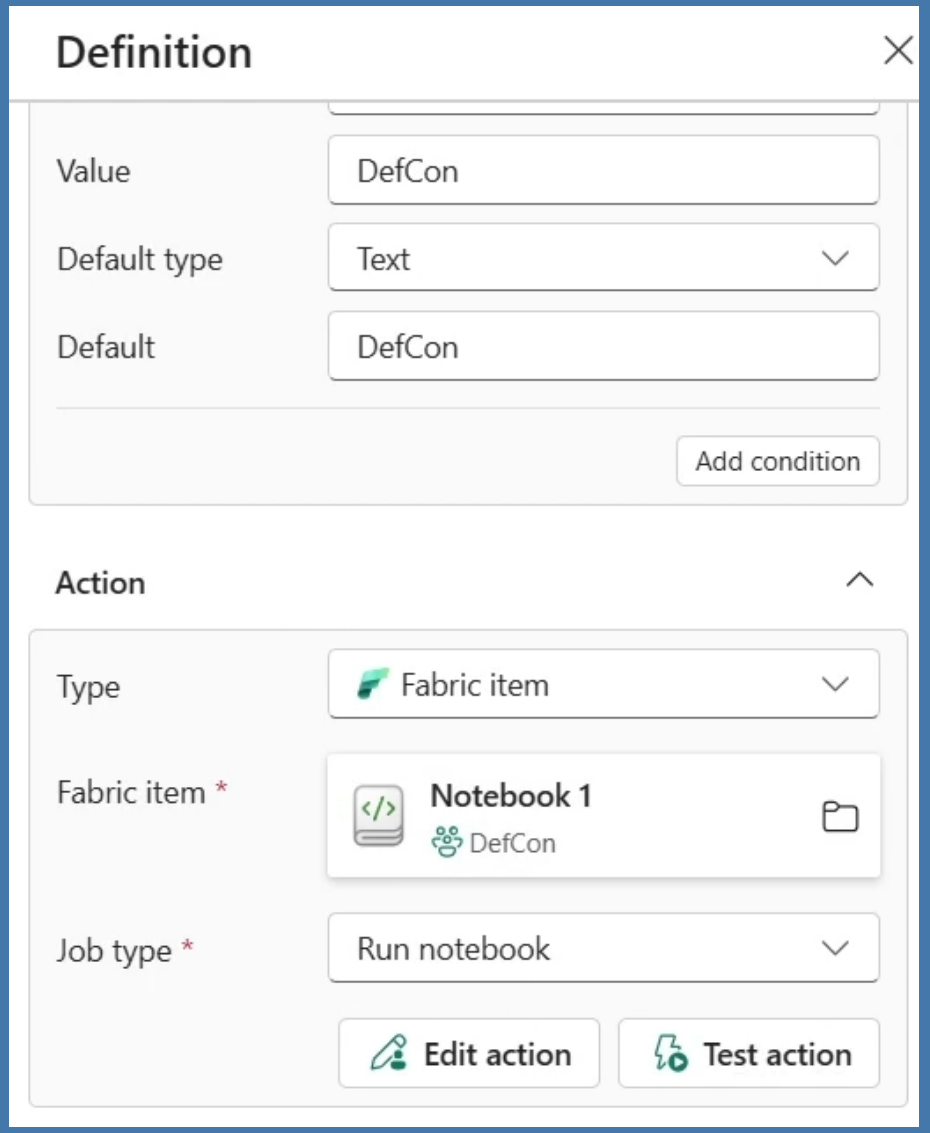
Definition
Value
DefCon
Default type
Text
Default
DefCon
Add condition
Action
Type
Fabric item
Fabric item *
Notebook 1
DefCon
Job type *
Run notebook
Edit action
Test action

Backdoor

- Notebook
- Azure Python SDK
- 3rd Party Library

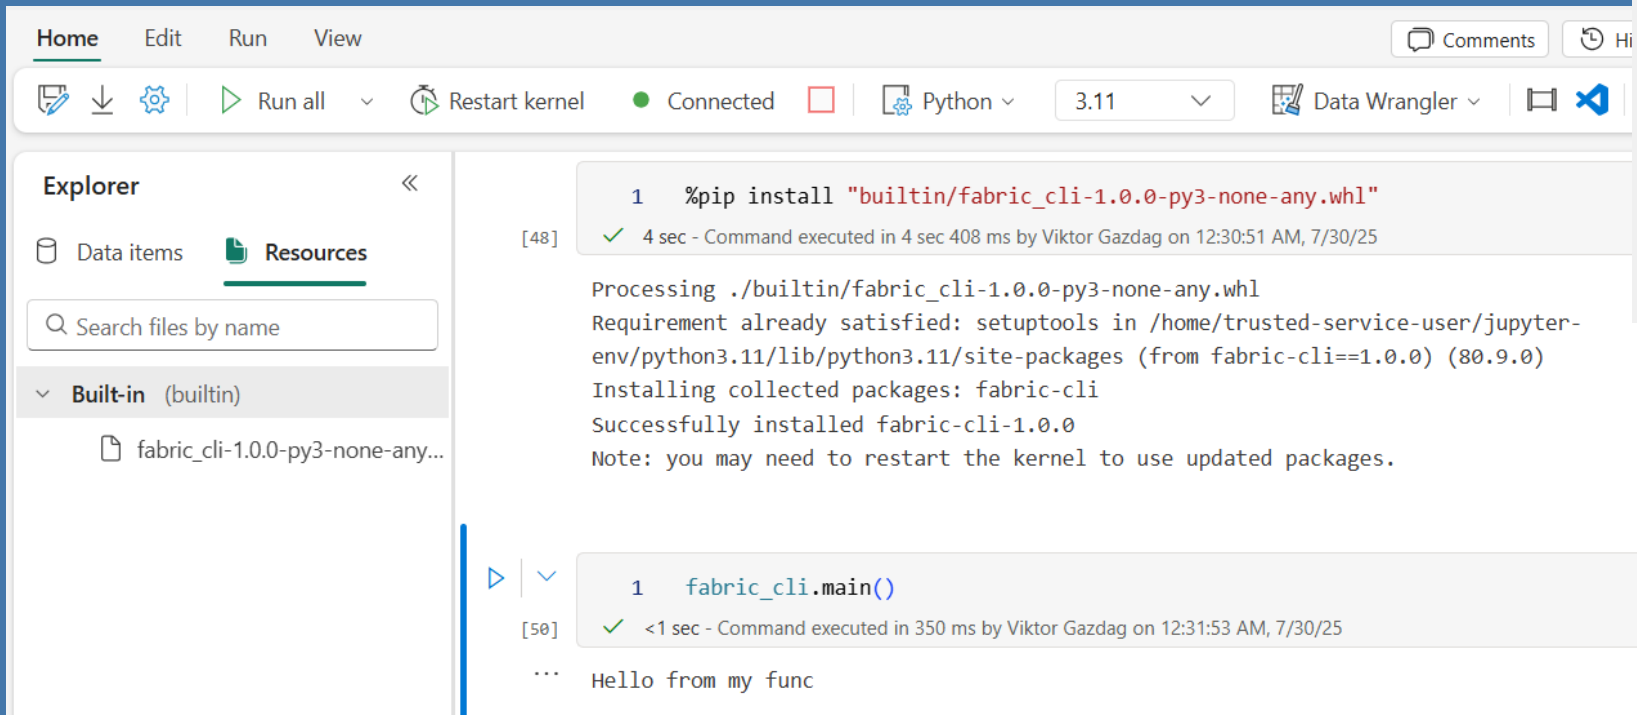

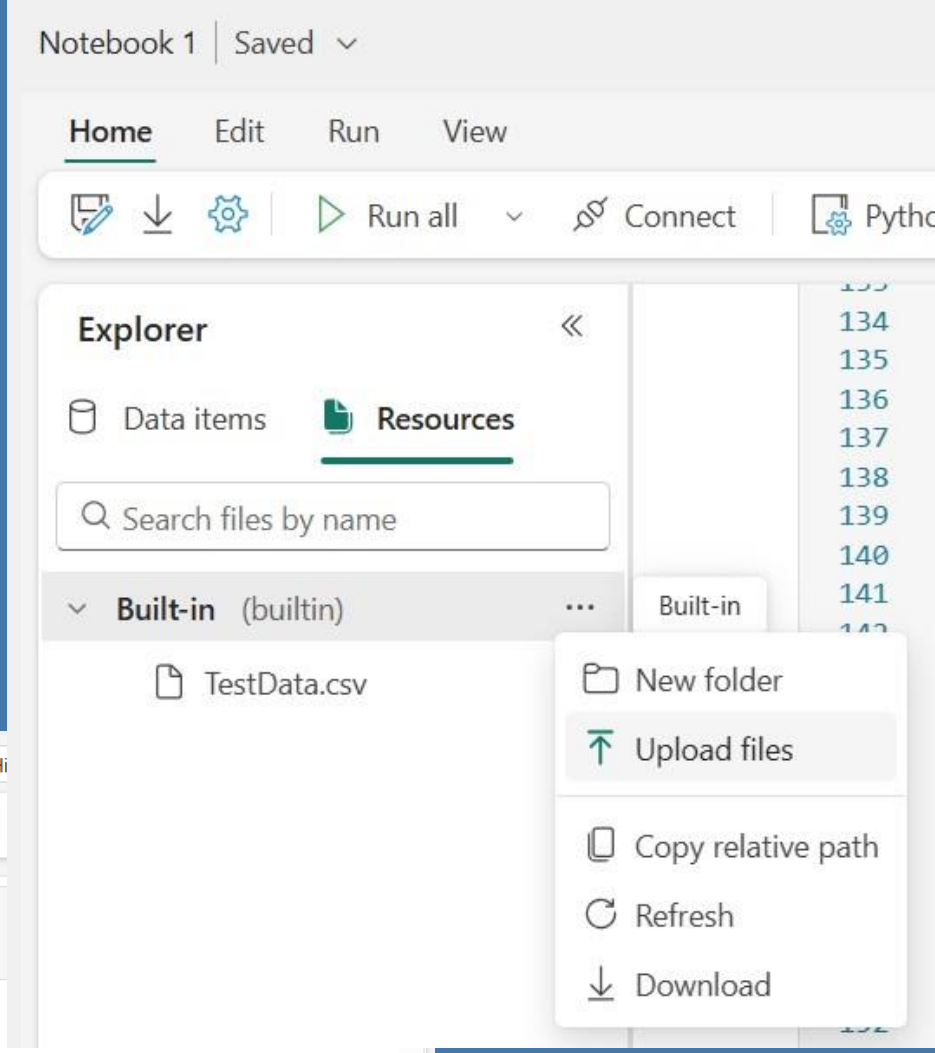

Azure Python SDK – Packages Install

```
1   !pip3 install azure-mgmt-msi
2   !pip3 install azure-mgmt-authorization
3   !pip3 install msrest
4   !pip3 install msrestazure
5   !pip3 install azure-mgmt-common
6   !pip3 install azure-mgmt-nspkg
7   !pip3 install azure-mgmt-compute
8   !pip3 install azure-mgmt-resource
9   !pip3 install azure-mgmt-network
10  !pip3 install azure-identity
```

Azure Python SDK - Import

```
#!/usr/bin/env python

import os
import random
import string
from datetime import datetime, timedelta

from azure.identity import ClientSecretCredential
from azure.mgmt.msi import ManagedServiceIdentityClient
from azure.mgmt.resource import ResourceManagementClient
from azure.mgmt.compute import ComputeManagementClient
from azure.mgmt.network import NetworkManagementClient
from azure.mgmt.authorization import AuthorizationManagementClient
from azure.mgmt.authorization.models import RoleAssignmentCreateParameters
from azure.core.exceptions import HttpResponseError
```

Azure Python SDK - Login

```
def main():
    # Create client credential
    credential = ClientSecretCredential(
        tenant_id=tenant_id,
        client_id=client_id,
        client_secret=client_secret
    )

    # Create clients
    resource_client = ResourceManagementClient(credential, subscription_id)
    msi_client = ManagedServiceIdentityClient(credential, subscription_id)
    compute_client = ComputeManagementClient(credential, subscription_id)
    network_client = NetworkManagementClient(credential, subscription_id)
    auth_client = AuthorizationManagementClient(credential, subscription_id)
```

Azure Python SDK – Role Assignment

```python
# Create managed identity
identity = create_managed_identity(msi_client, resource_group_name, managed_identity_name, location)
print(f"Created managed identity: {identity.name} with ID: {identity.id}")

# Assign Contributor role to the managed identity
role_assignment = assign_role(
    auth_client,
    identity.principal_id,
    contributor_role_id,
    f"/subscriptions/{subscription_id}/resourceGroups/{resource_group_name}"
)
print(f"Assigned Contributor role to managed identity")
```

Azure Python SDK – Role Assignment 2

```python
def assign_role(auth_client, principal_id, role_definition_id, scope):
    """Assign role to the managed identity."""
    # Generate a random UUID for the role assignment name
    role_assignment_name = f"{principal_id}-{role_definition_id}"[:36]

    role_assignment_params = RoleAssignmentCreateParameters(
        role_definition_id=f"/subscriptions/{subscription_id}/providers/Microsoft.Authorization/roleDefinitions/{role_definition_id}",
        principal_id=principal_id,
        principal_type="ServicePrincipal"  # Specify principal type to avoid replication delay issues
    )

    try:
        return auth_client.role_assignments.create(
            scope=scope,
            role_assignment_name=role_assignment_name,
            parameters=role_assignment_params
        )
    except HttpResponseError as e:
        # If role assignment already exists, continue
        if "already exists" in str(e):
            print(f"Role assignment already exists for principal {principal_id}")
            return None
        raise
```

Azure Python SDK – Managed Identity

```python
def create_managed_identity(msi_client, resource_group_name, identity_name, location):
    """Create a user-assigned managed identity."""
    try:
        return msi_client.user_assigned_identities.get(resource_group_name, identity_name)
    except HttpResponseError:
        # Create the managed identity
        identity = msi_client.user_assigned_identities.create_or_update(
            resource_group_name,
            identity_name,
            {"location": location}
        )

        # Add a small delay to allow for replication
        print("Waiting for managed identity to propagate...")
        import time
        time.sleep(30)  # Wait 30 seconds for replication

        return identity
```

Backdoor

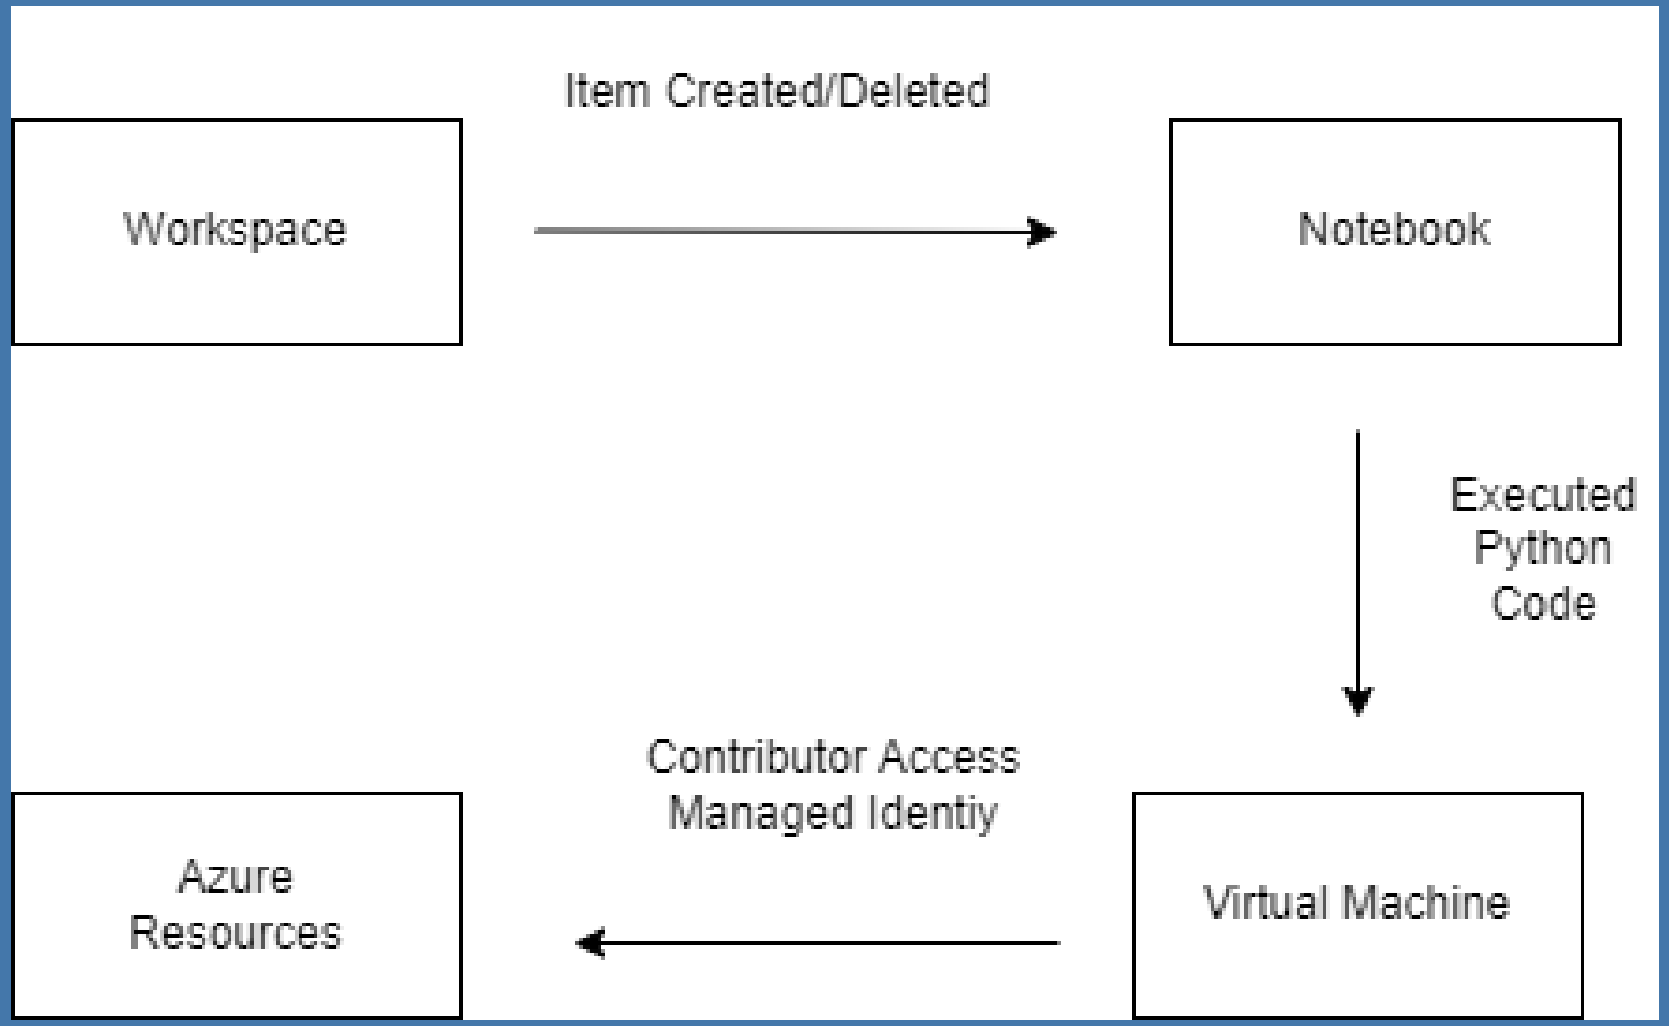
Item Created/Deleted
Workspace
Notebook
Executed
Python
Code
Contributor Access
Managed Identiy
Azure
Resources
Virtual Machine

Backdoor

- Demo video – 7:42 mins

Data Exfiltration

- Data Pipeline
- Notebook
- Shortcut
- SQL Endpoint
- Mirrored DB

Data Exfiltration

- Data Pipeline
- Add A Step To Copy Data To Somewhere Else
- Use Notebook With Hidden Code

Data Exfiltration

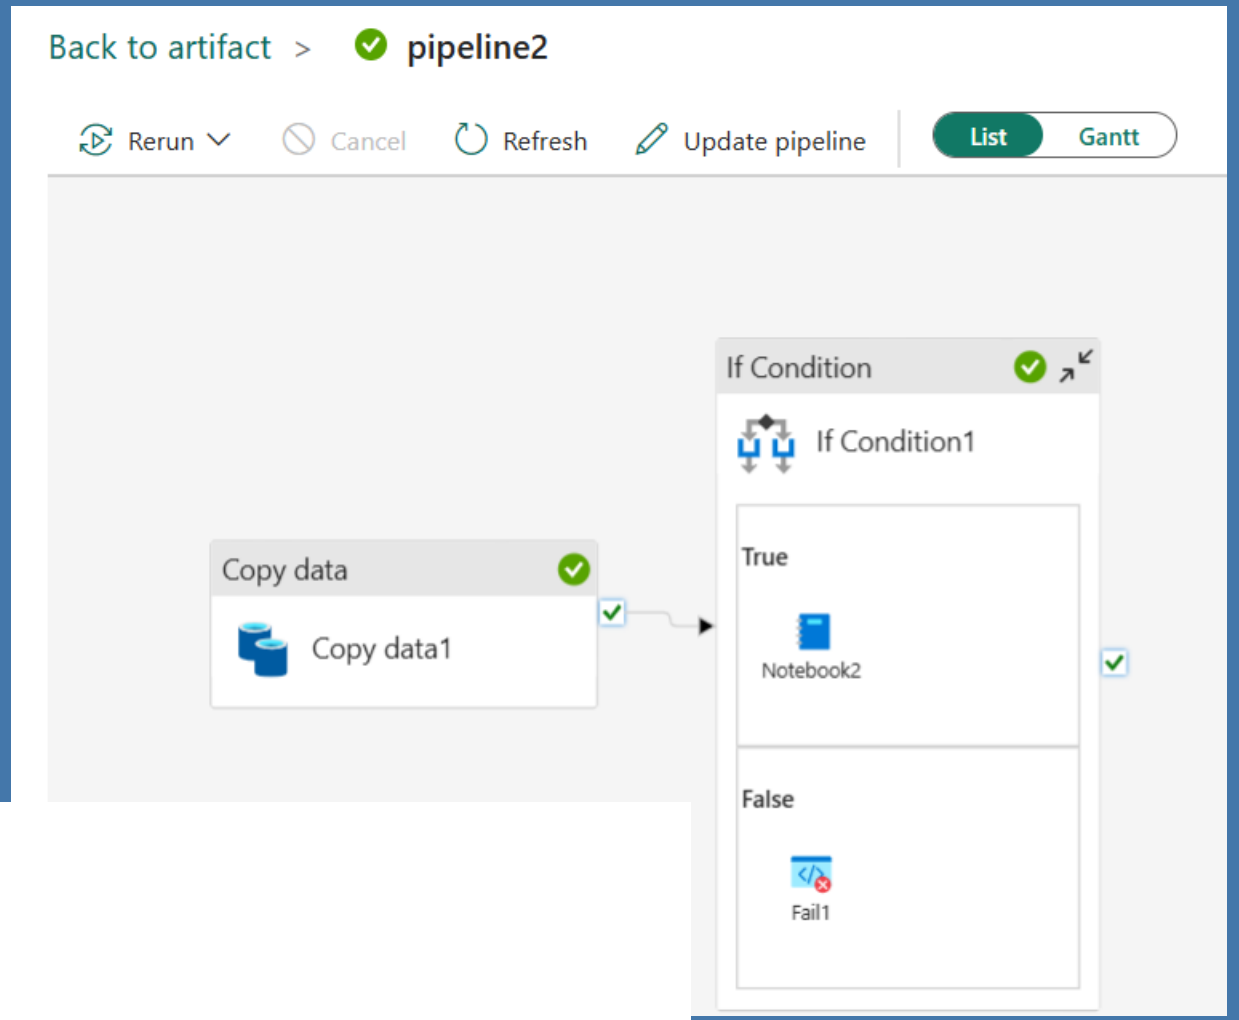

Activity runs

Filter by keyword

Showing 3 items

Activity name	Activity sta...	Run start	Duration	Input	Output
Copy data1	Succeeded	7/29/2025, 12:30:28 AM	20s		
If Condition1	Succeeded	7/29/2025, 12:30:49 AM	25s		
Notebook2	Succeeded	7/29/2025, 12:30:49 AM	23s		

Data Exfiltration

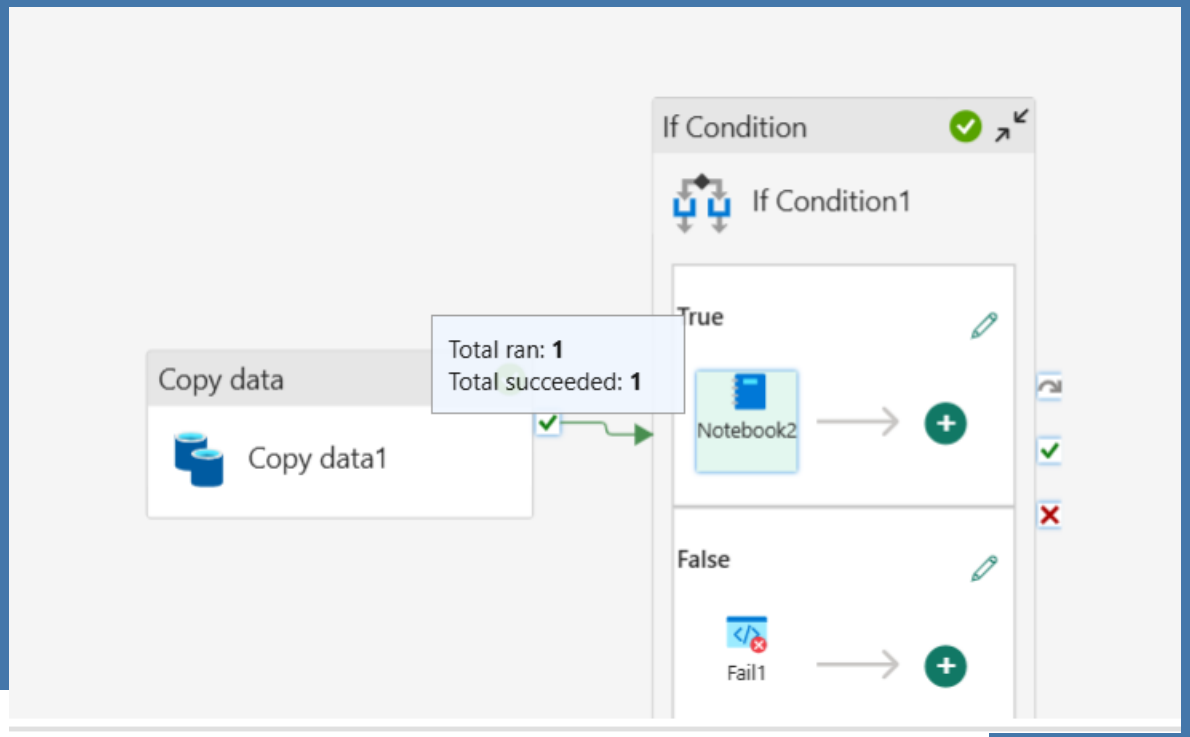

Parameters | Variables | Settings | **Output** | Library variables (preview)

Pipeline run ID 4a595ba7-d68d-48e4-93b5-9f64289d0dbc View run detail **Pipeline status** Succeeded

Filter by keyword Showing 3 items

Activity name	Activity st...	Run start	Duration	Input	Output
Copy data1	Succeeded	7/29/2025, 12:30:28 AM	20s		
If Condition1	Succeeded	7/29/2025, 12:30:49 AM	25s		
Notebook2	Succeeded	7/29/2025, 12:30:49 AM	23s		

Data Exfiltration

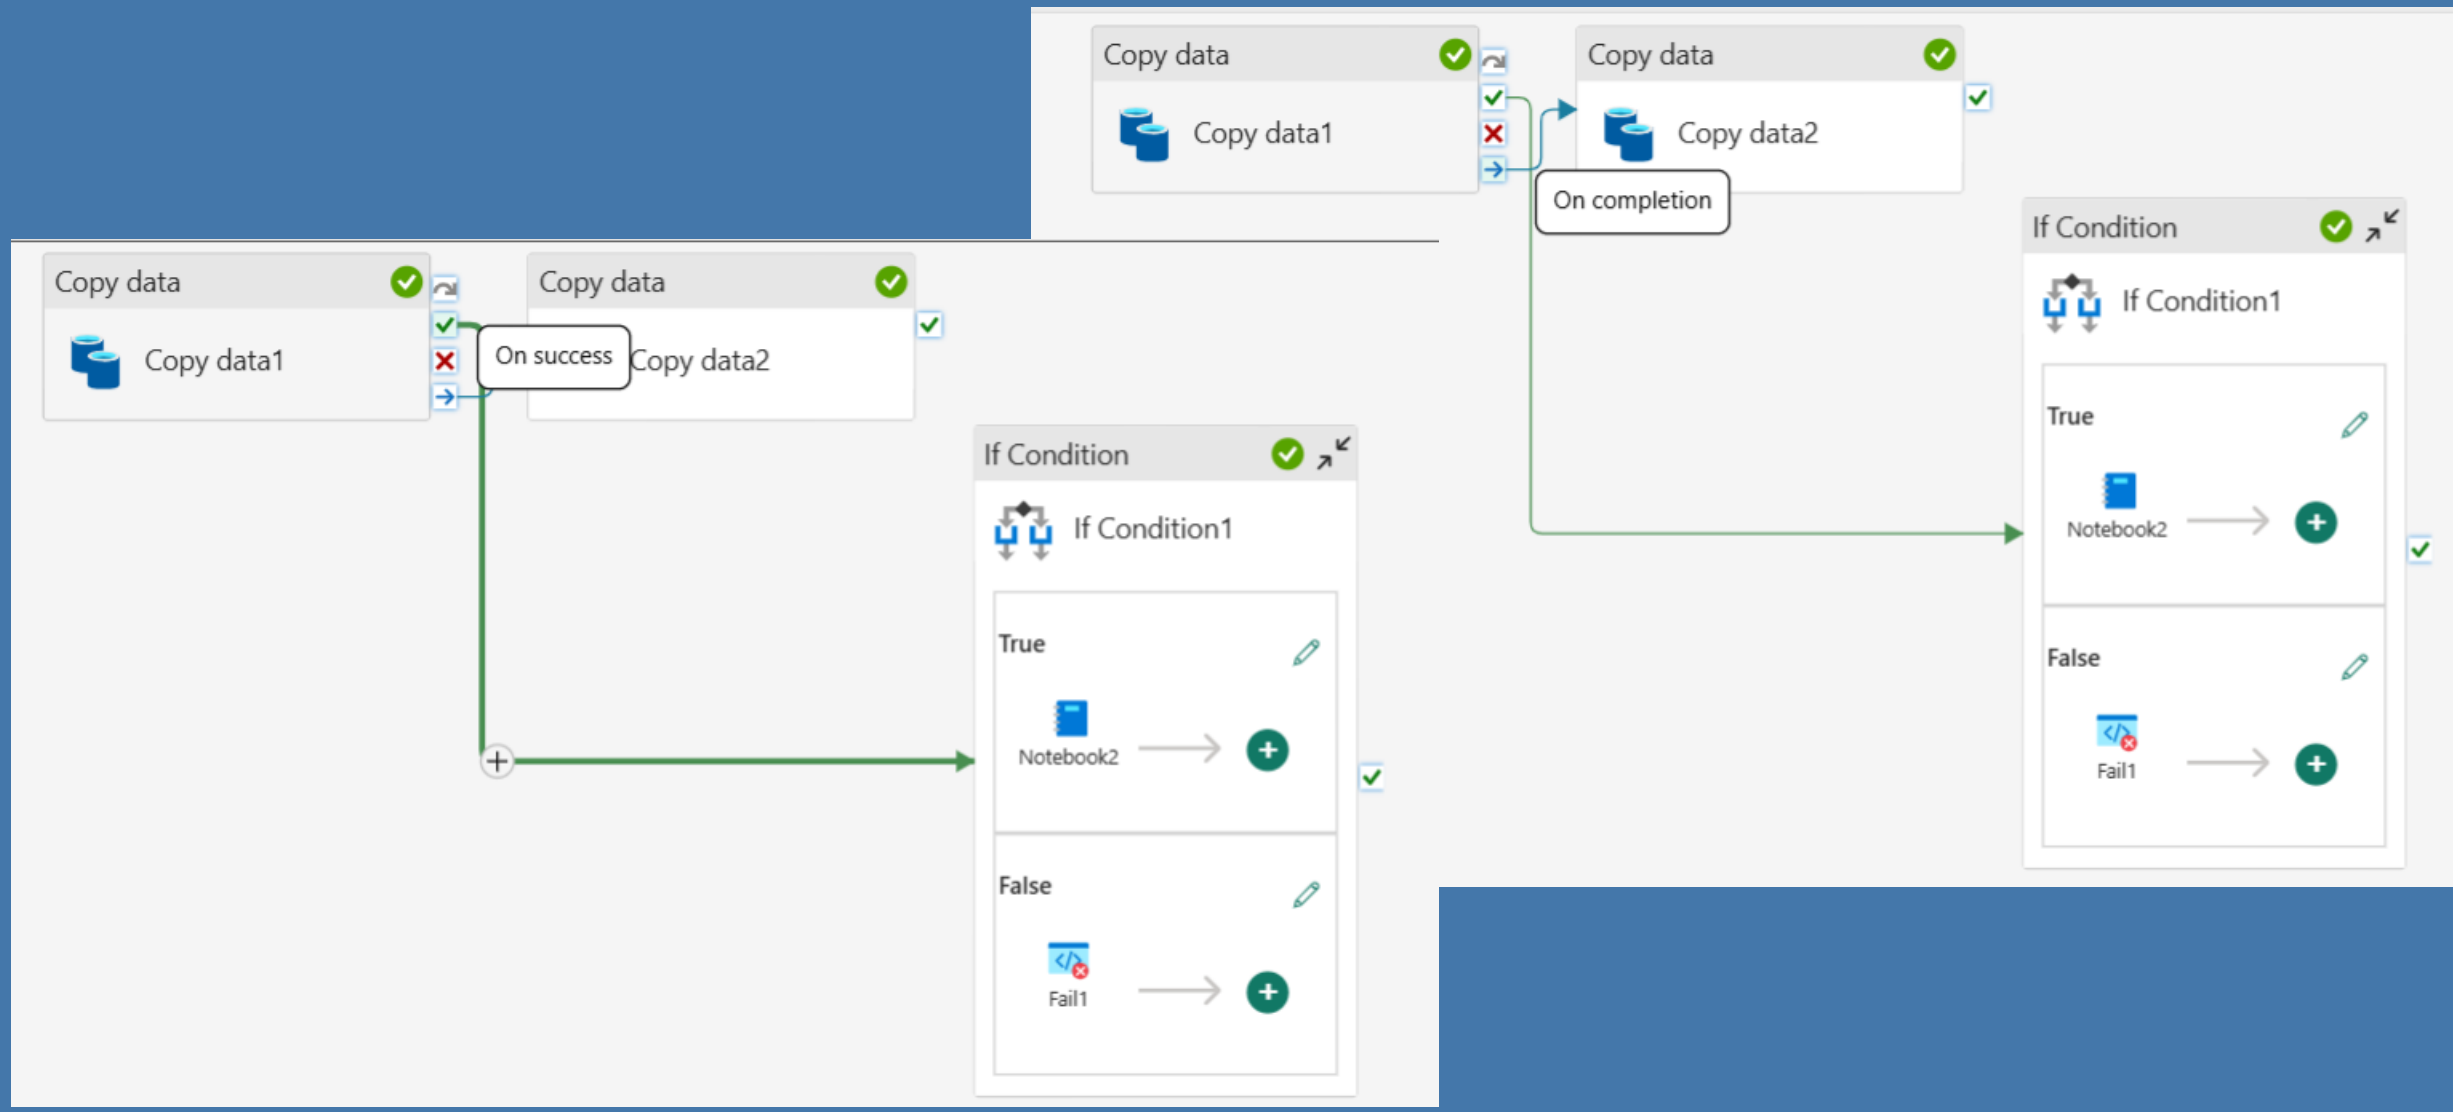

Data Exfiltration

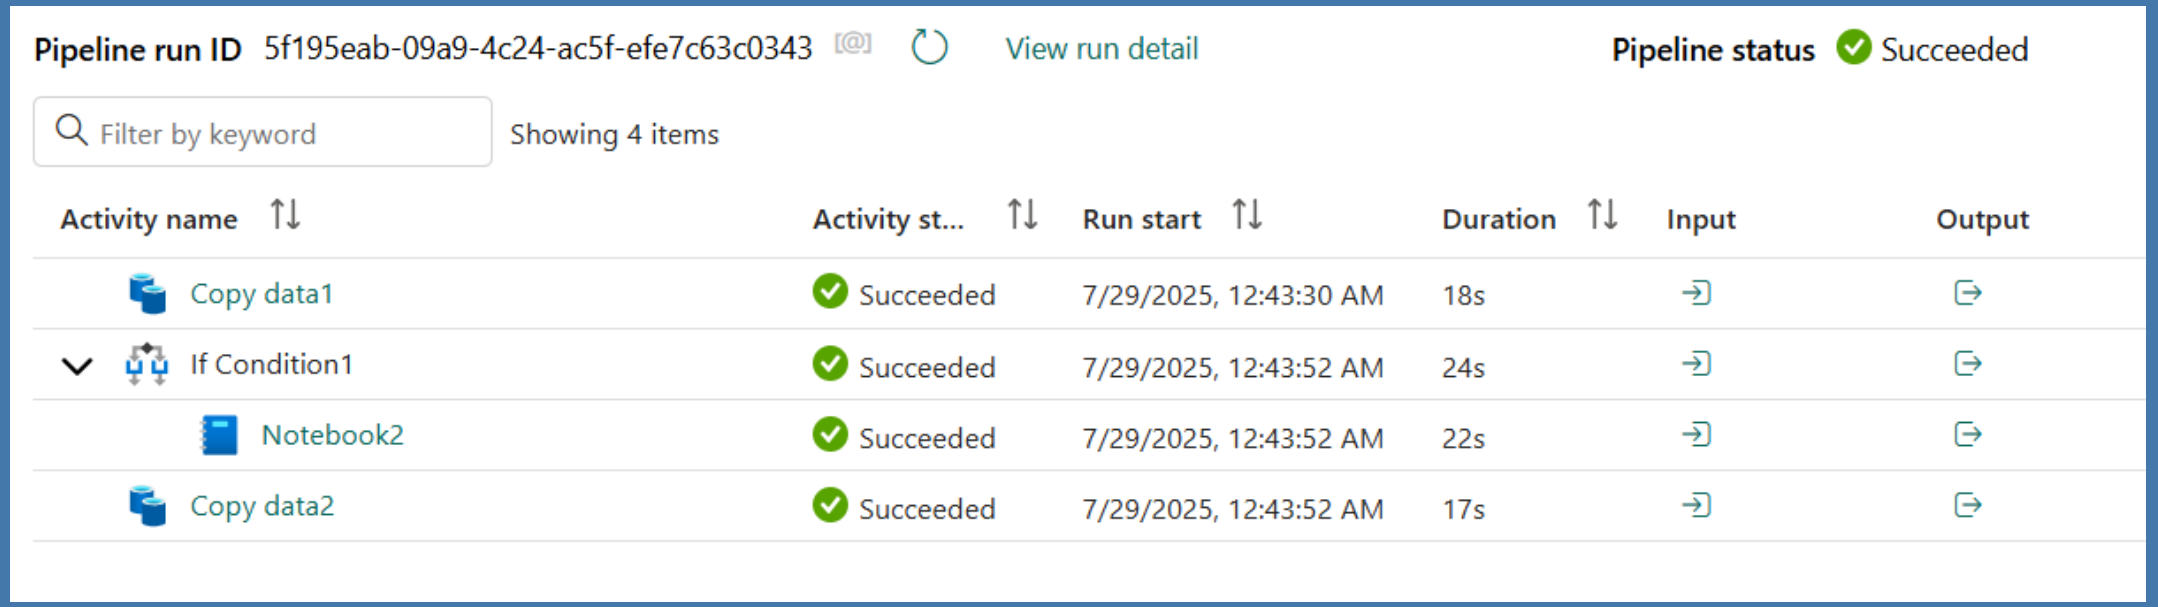

Data Exfiltration

- Notebook
- Python Runtime
- %pip% Install Custom Libraries
  - Built-in Resource
  - Spark Environment Custom Library
- Internet Access Unless – Private Links For Secure Access To Fabric (Can Break Things)
- Scheduled Notebook Running

Data Exfiltration

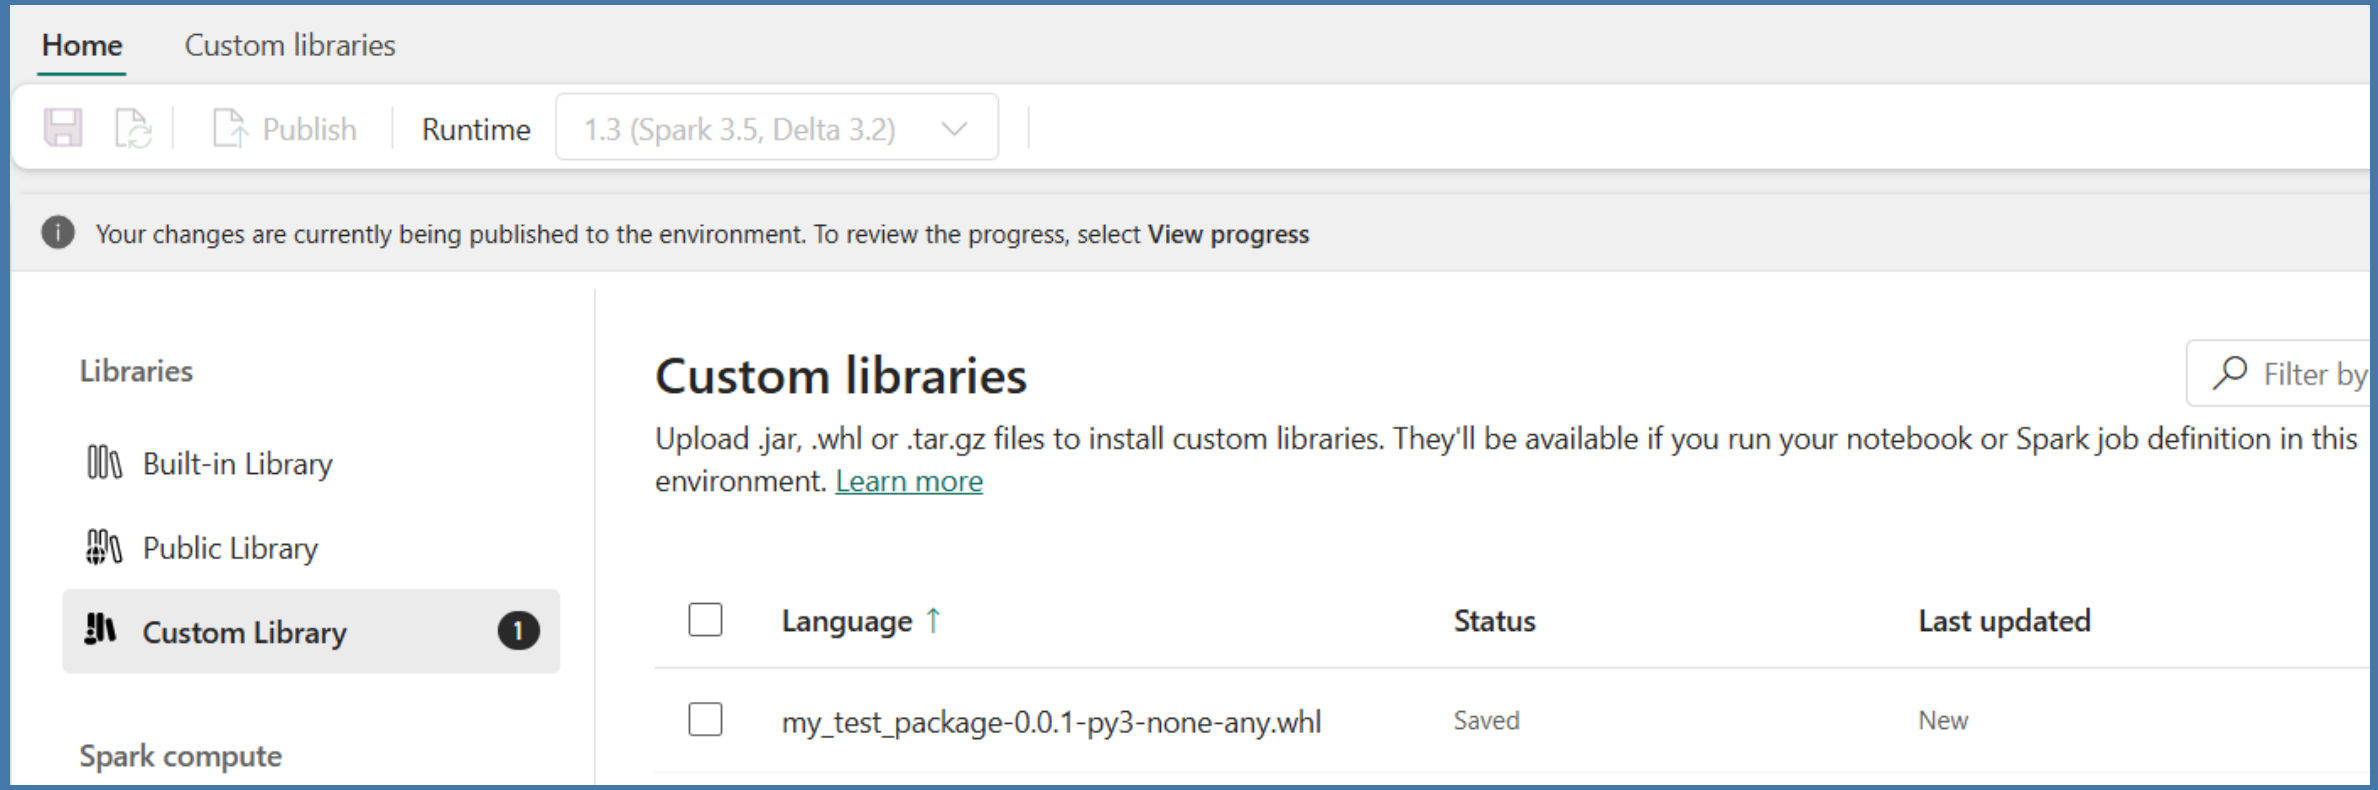

Data Exfiltration

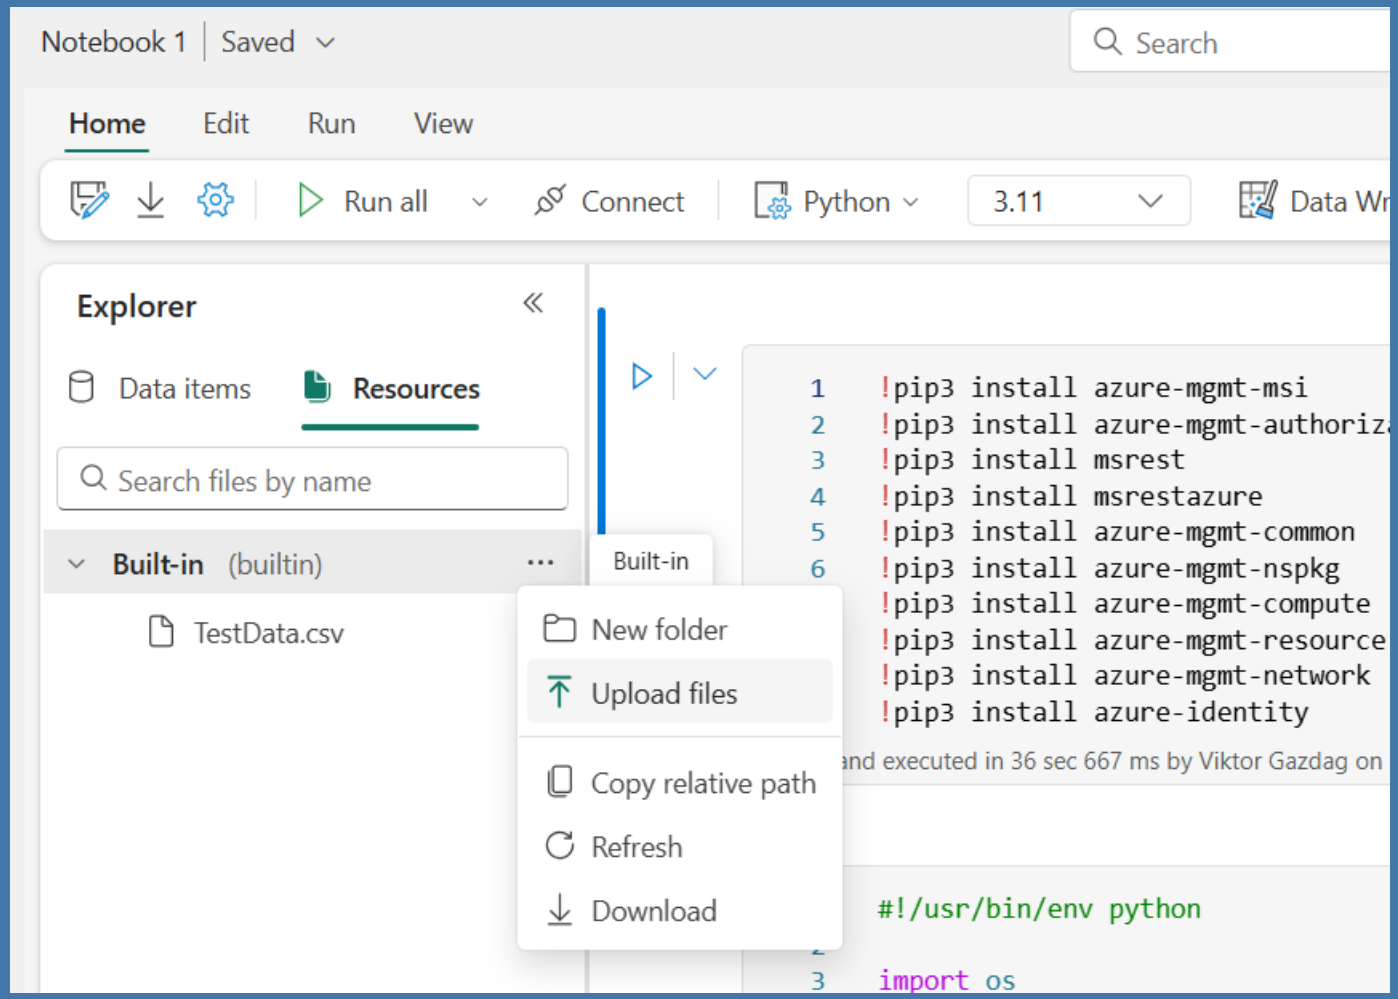

Data Exfiltration

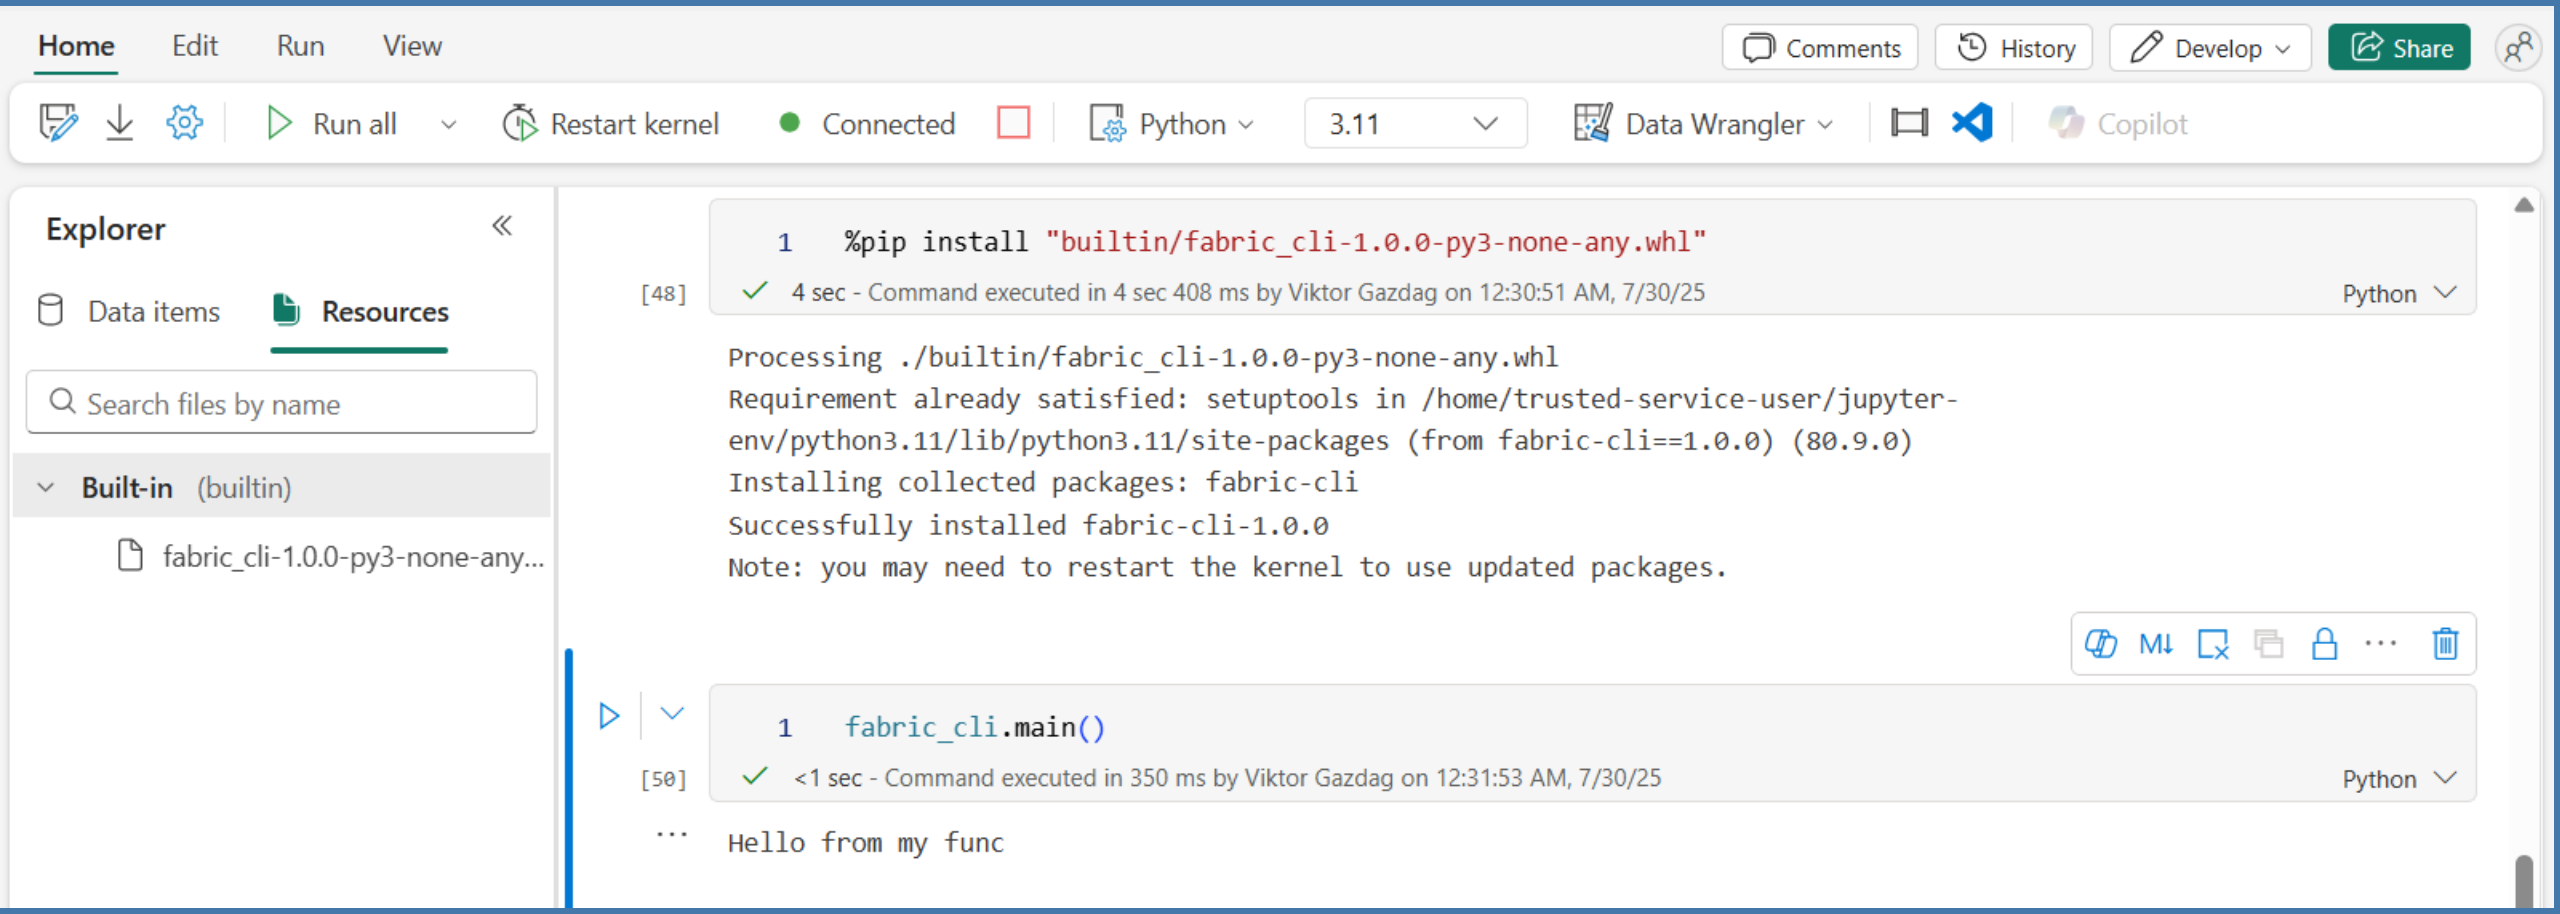

Data Exfiltration

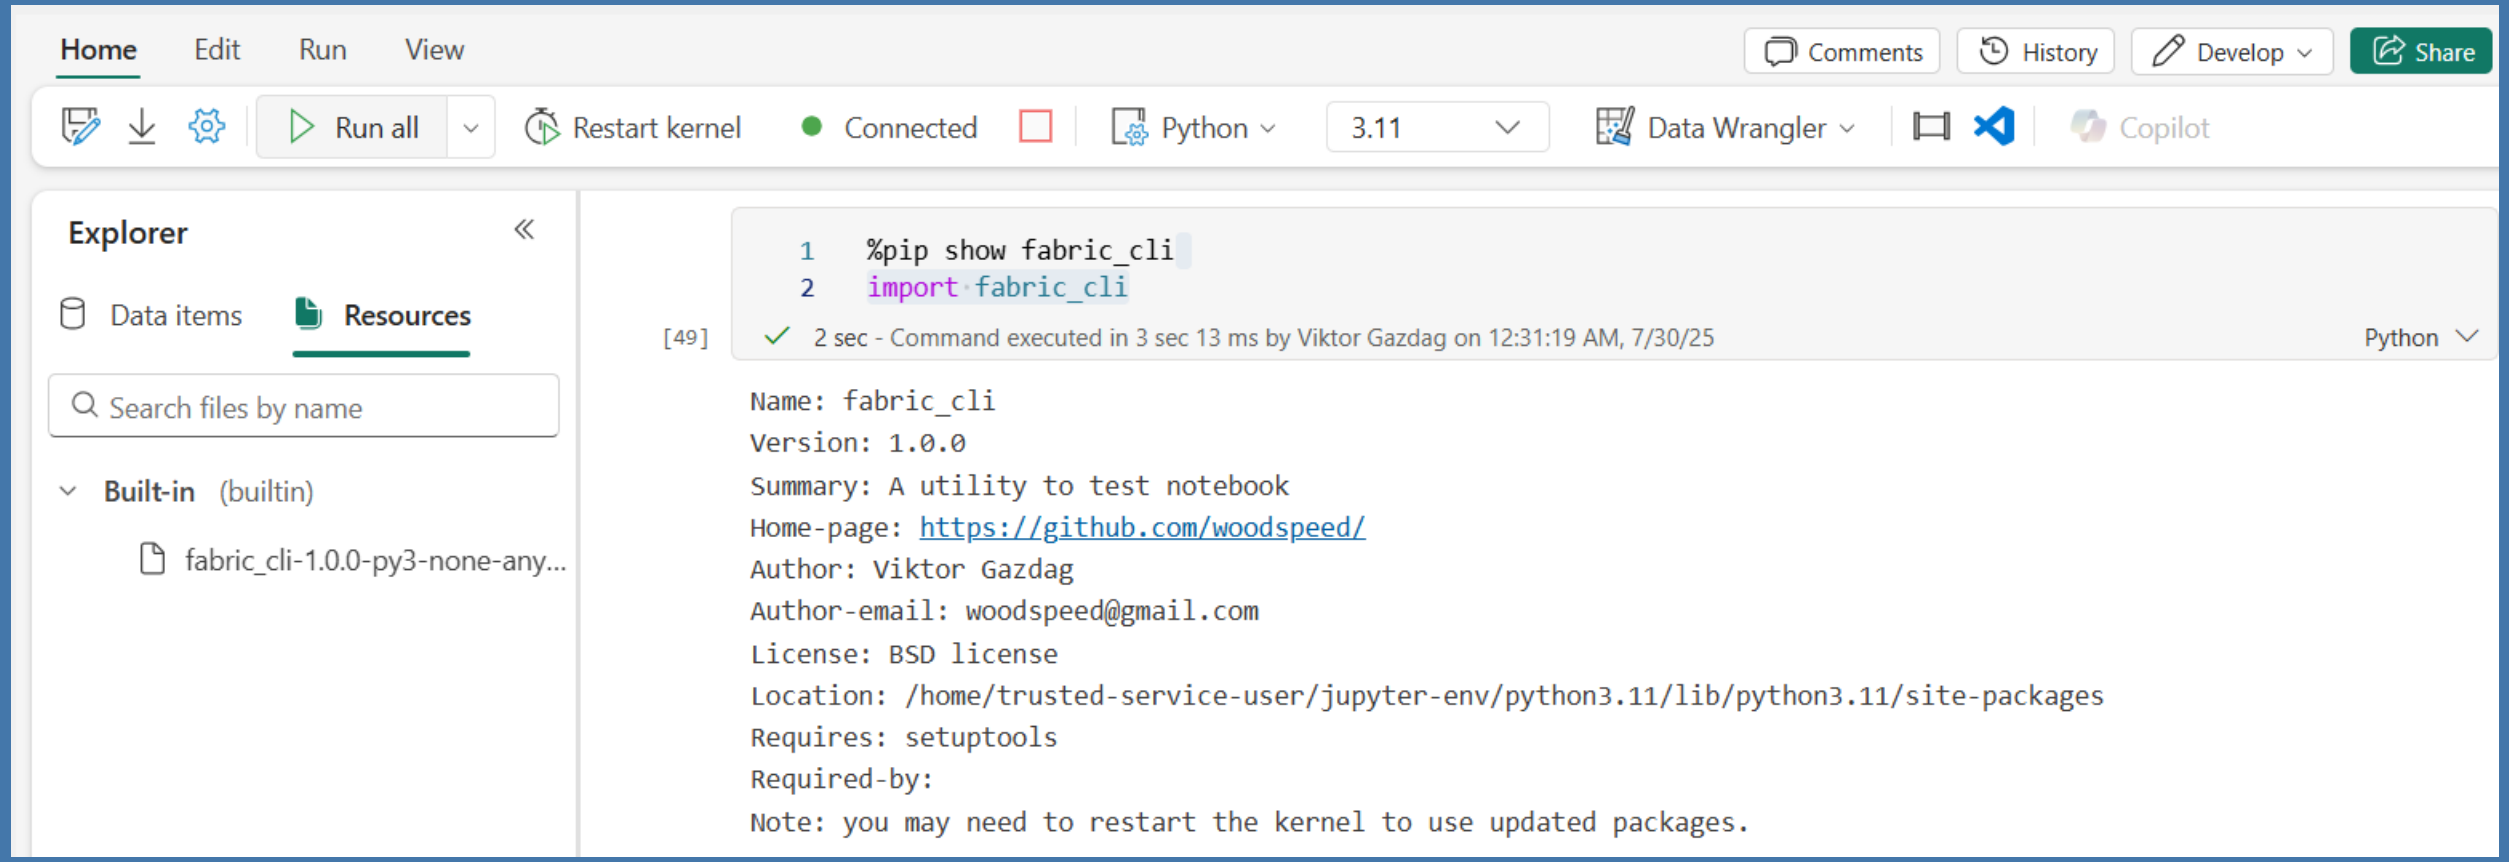

Data Exfiltration

- Shortcut (Symbolic Link)
  - Point To Target
  - Work On Target DB
- Internal To Fabric
  - Lakehouse
  - Warehouse
- External To Fabric
  - Datalake
  - S3
- Require ReadAll On Source And Write+ReadAll On Target For Reading Data

Data Exfiltration

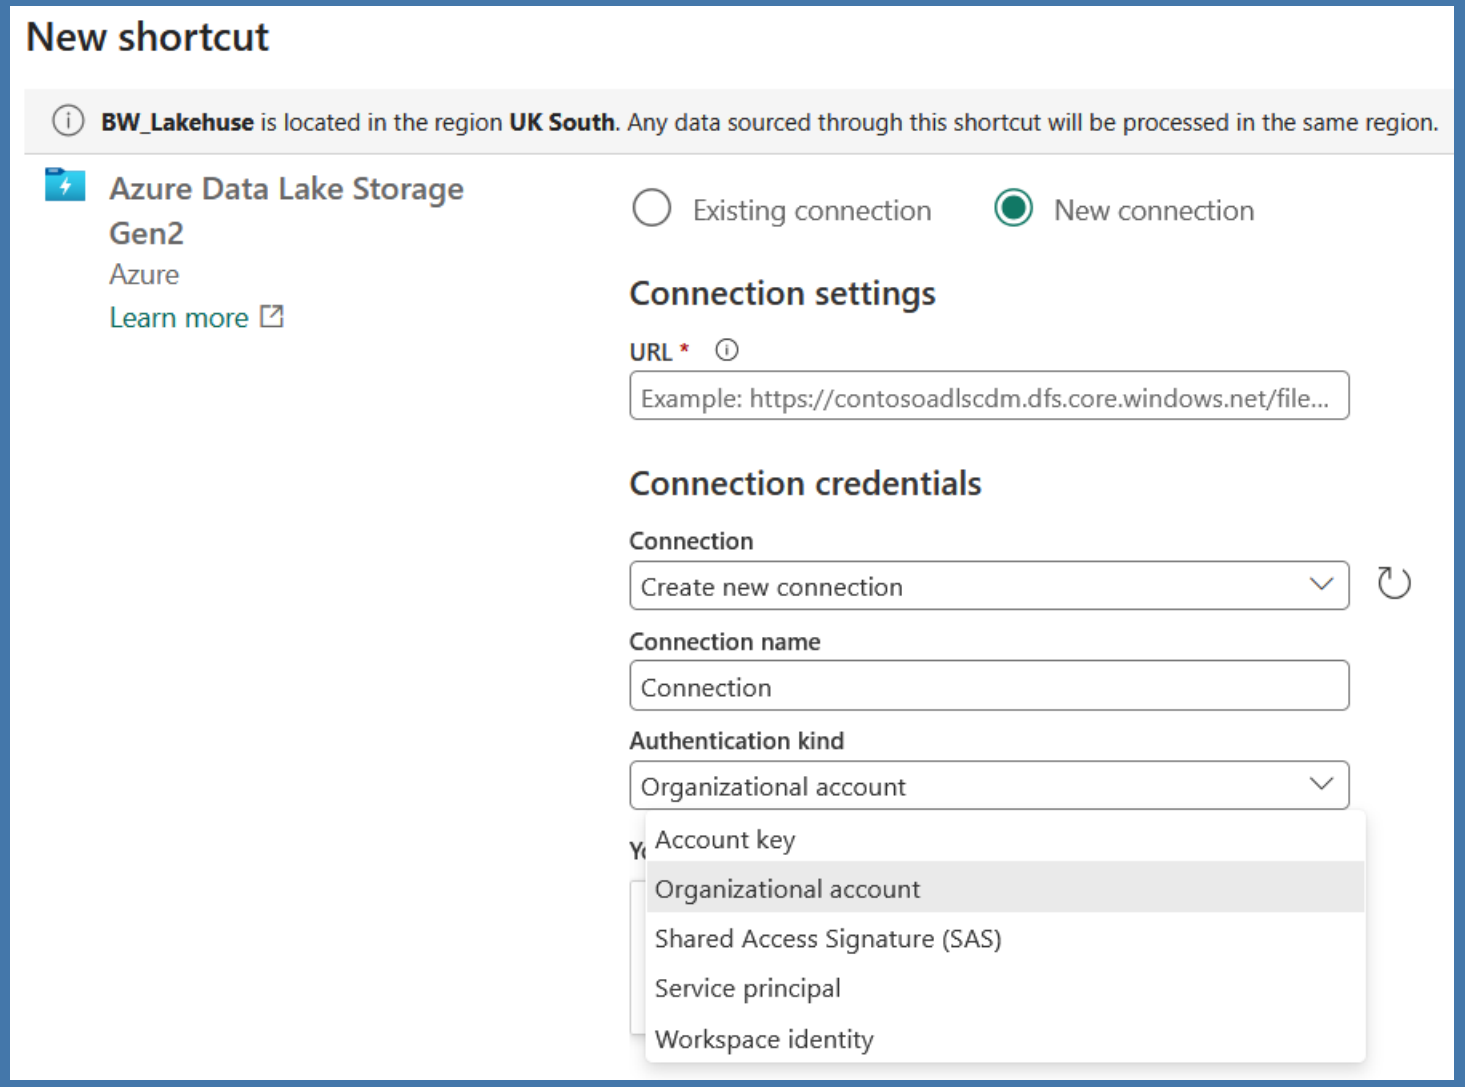

Data Exfiltration

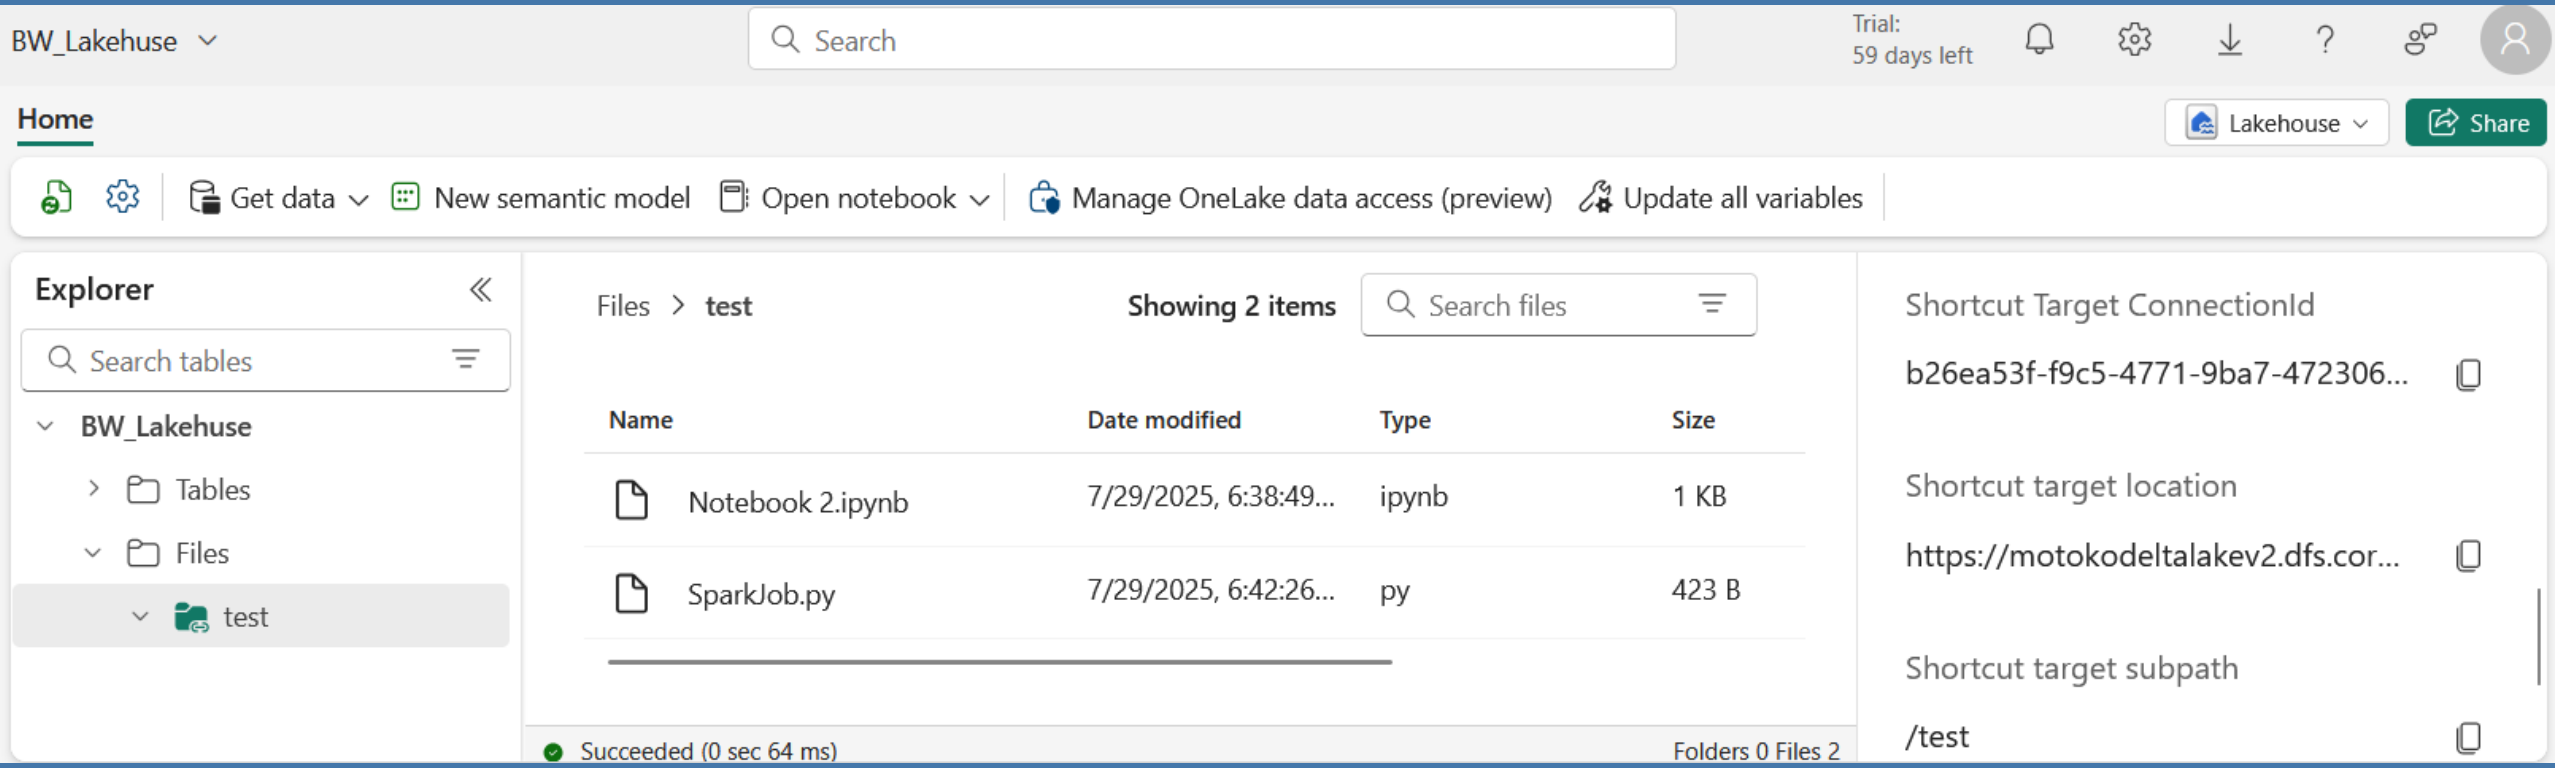

Data Exfiltration

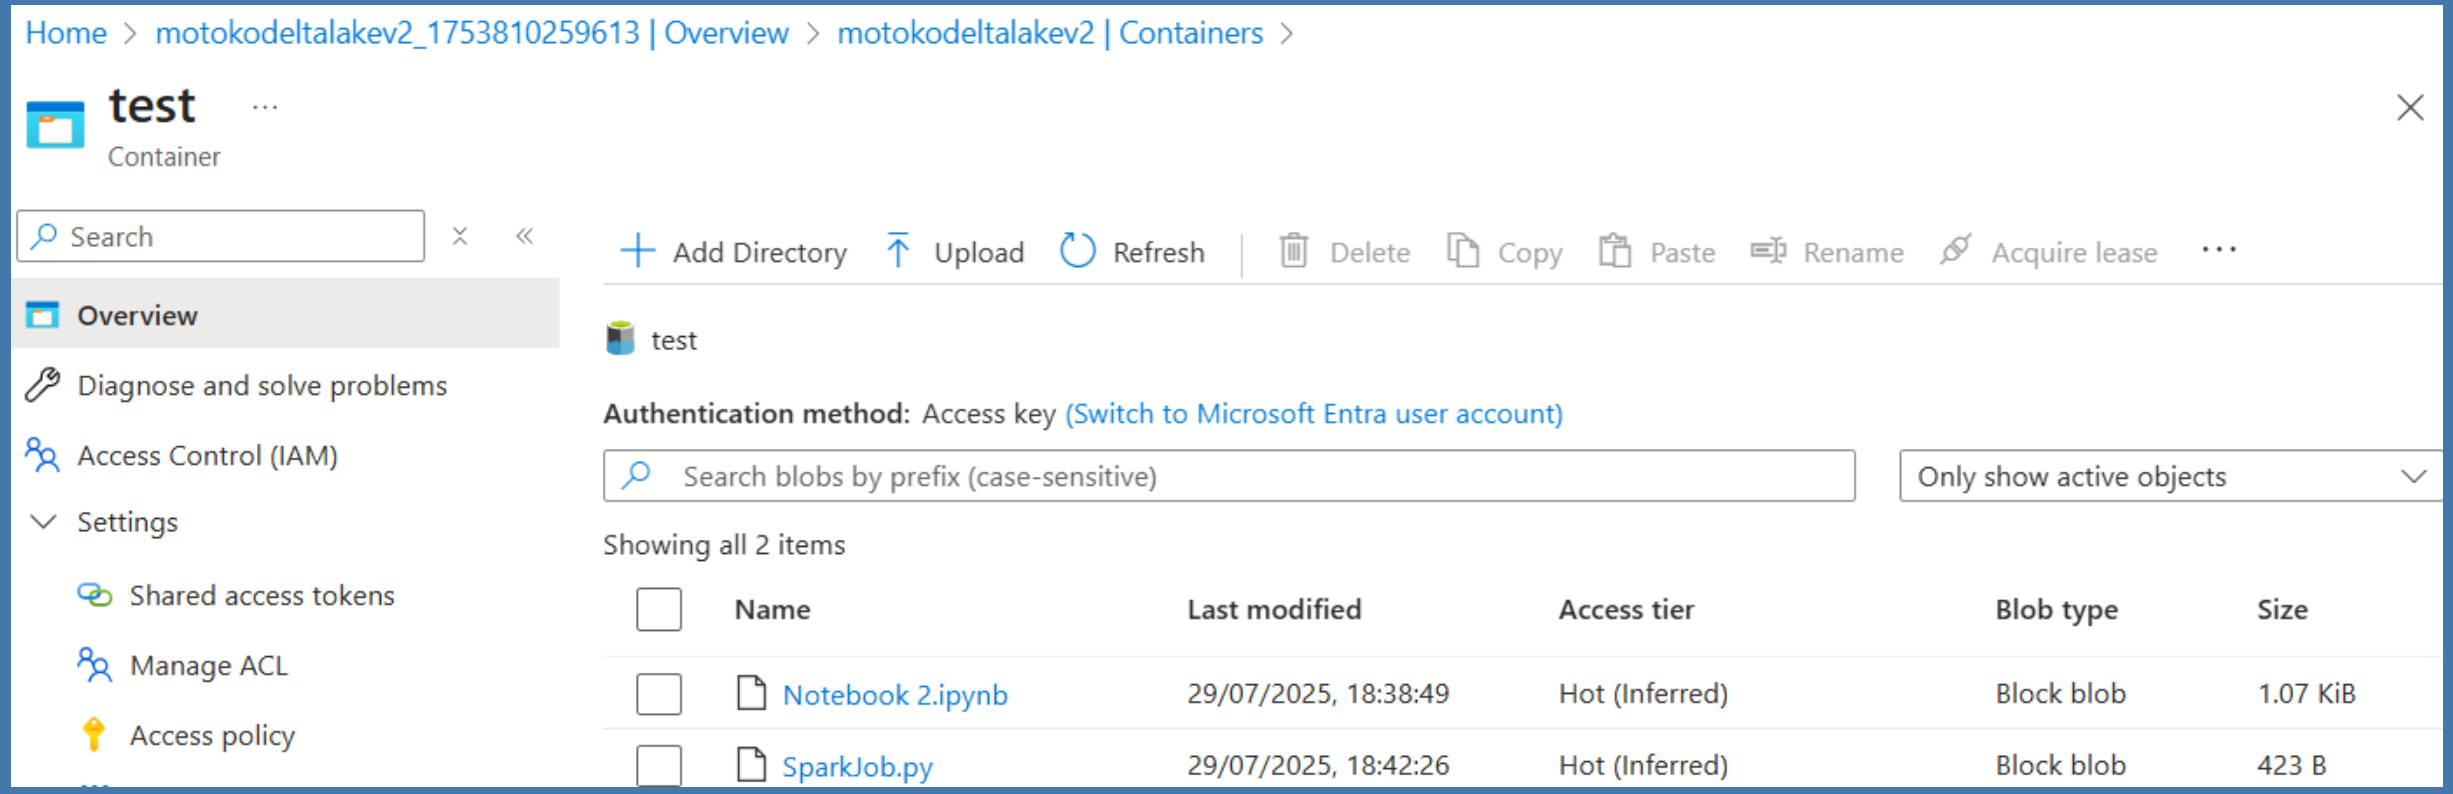

Data Exfiltration

- SQL Endpoint
- Automatically Created With Lakehouse
- Read-only
- Share (Read (able to connect) + Permission At Database Level = “Hidden Access”
- Share With Entra ID Guest Account And Grant Select On Database

CloudLake

Add user

woodspeed

Direct access

People and groups with access	Email Address	Role	Permissions
woodspeed_blackwombat.com...	woodspeed_blackwombat.com#EXT#@mot...		Read

File Edit View Query Git Project Tools Extensions Window Help Search Solution1 Sign in

New Query

CloudLake Execute

Object Explorer

Connect

krkk3naiph3evi6bqybgailivi-7lwrtzjngyxedgwpilj3xmk5jm...
- Databases
  - System Databases
  - CloudLake
    - Tables
      - System Tables
      - External Tables
      - dbo.testdata
      - Dropped Ledger Tables
    - Views
    - External Resources
    - Programmability
    - XEvent Profiler
    - Security
- Security

SQLQuery1.s...at.com (62))*

```
1  SELECT TOP (1000) [Name]
2        ,[CyberDeck]
3    FROM [dbo].[testdata]
4
5    select user
```

100 % No issues found Ln: 5

Results Messages

	Name	CyberDeck
1	V	No
2	Viktor	Yes
3	Songbird	Yes

	(No column name)
1	woodspeed@blackwombat.com

Data Exfiltration

	principal_name	permission_name	state_desc	table_name
1	woodspeed@blackwombat.com	SELECT	GRANT	testdata
2	woodspeed@blackwombat.com	SELECT	GRANT	testdata

GRANT select ([Name], [CyberDeck])
ON [CloudLake].[dbo].[testdata]
to “username"

Data Exfiltration

- Mirrored DB (Transaction Log)
- Replicate DB to Fabric
  - From Target Azure DB To Attacker Fabric OneLake
  - Work On Delta Parquet In Fabric
- Drops The Mirrored Table Data From Fabric OneLake

Data Exfiltration

- Replicate DB To Fabric
  - From Target Azure DB To Attacker Fabric OneLake
  - Work On Delta Parquet In Fabric

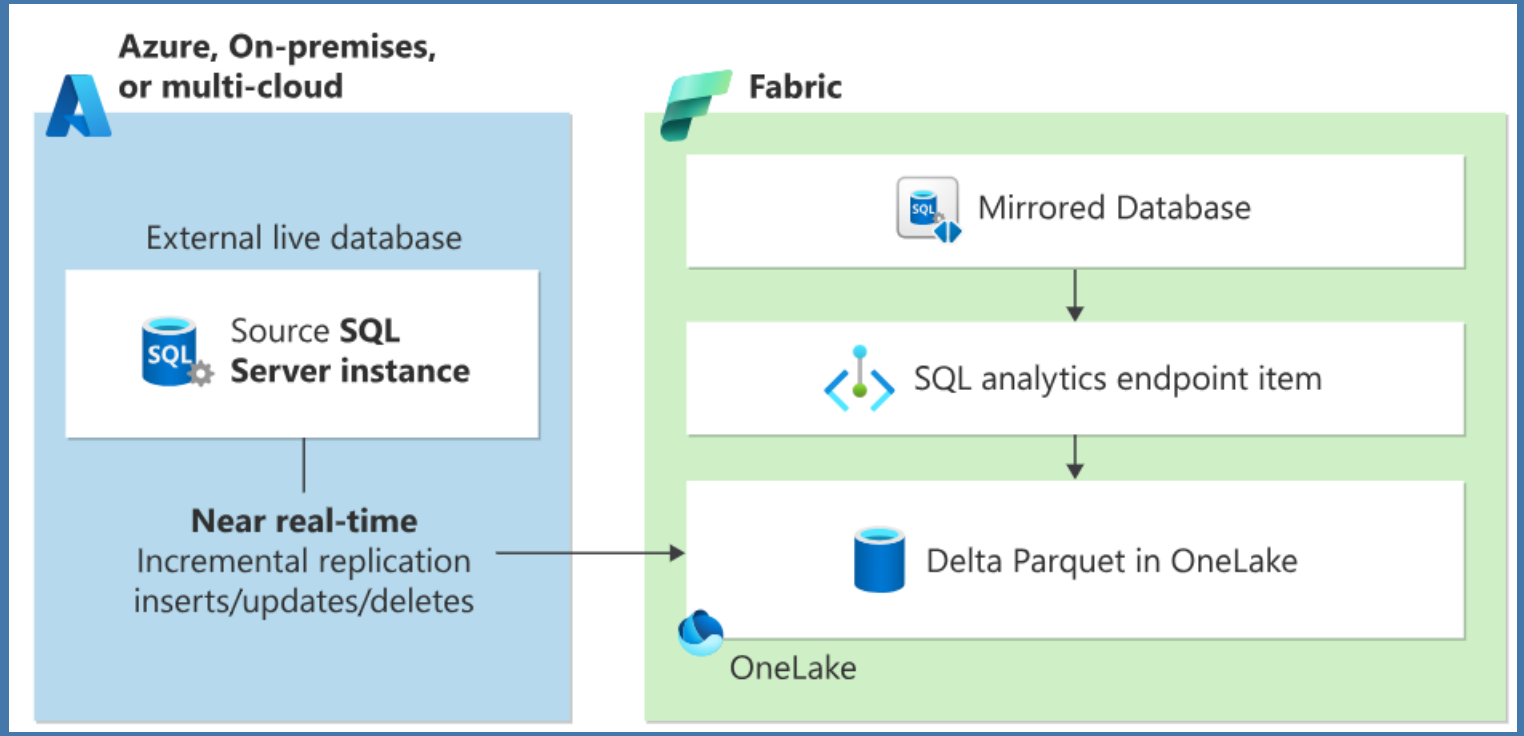

Summary

- Tenant Level Settings
  - Sharing
  - Service Principal
  - User Function
- Backdoor
  - Notebook
- Data Exfiltration
  - Pipeline
  - Shortcut
  - Notebook
  - SQL Endpoint
  - Mirrored DB

Links

- Azure Python SDK
  - https://learn.microsoft.com/en-gb/python/api/overview/azure/compute?view=azure-python
- Fabric Settings Tool
  - https://github.com/dcnsakthi/blogs/blob/main/fabric/admin/TenantSettings.md
- Fabric Tenant Settings Descriptions
  - https://learn.microsoft.com/en-us/fabric/admin/tenant-settings-index
- Microsoft Fabric Security Whitepaper
  - https://learn.microsoft.com/en-us/fabric/security/white-paper-landing-page

Questions & Contacts

- X (Twitter)
  - @wucpi
- Email
  - viktor.gazdag@nccgroup.com
- Conference Slides
  - https://github.com/woodspeed/conferences